Karlsruhe Institute of Technology

Clemens Lode, Urban Richter, Hartmut Schmeck
Karlsruhe Institute of Technology (Germany)
Institute AIFB
July, GECCO 2010

Adaption of XCS to Multi-Learner Predator/Prey Scenarios

http://www.flickr.com/photos/shreeram

Outline

http://www.flickr.com/photos/yathin

- Learning Classifier Systems
- XCS in Predator/Prey Scenarios
- Adapting the Reward Function
- Experimental Results

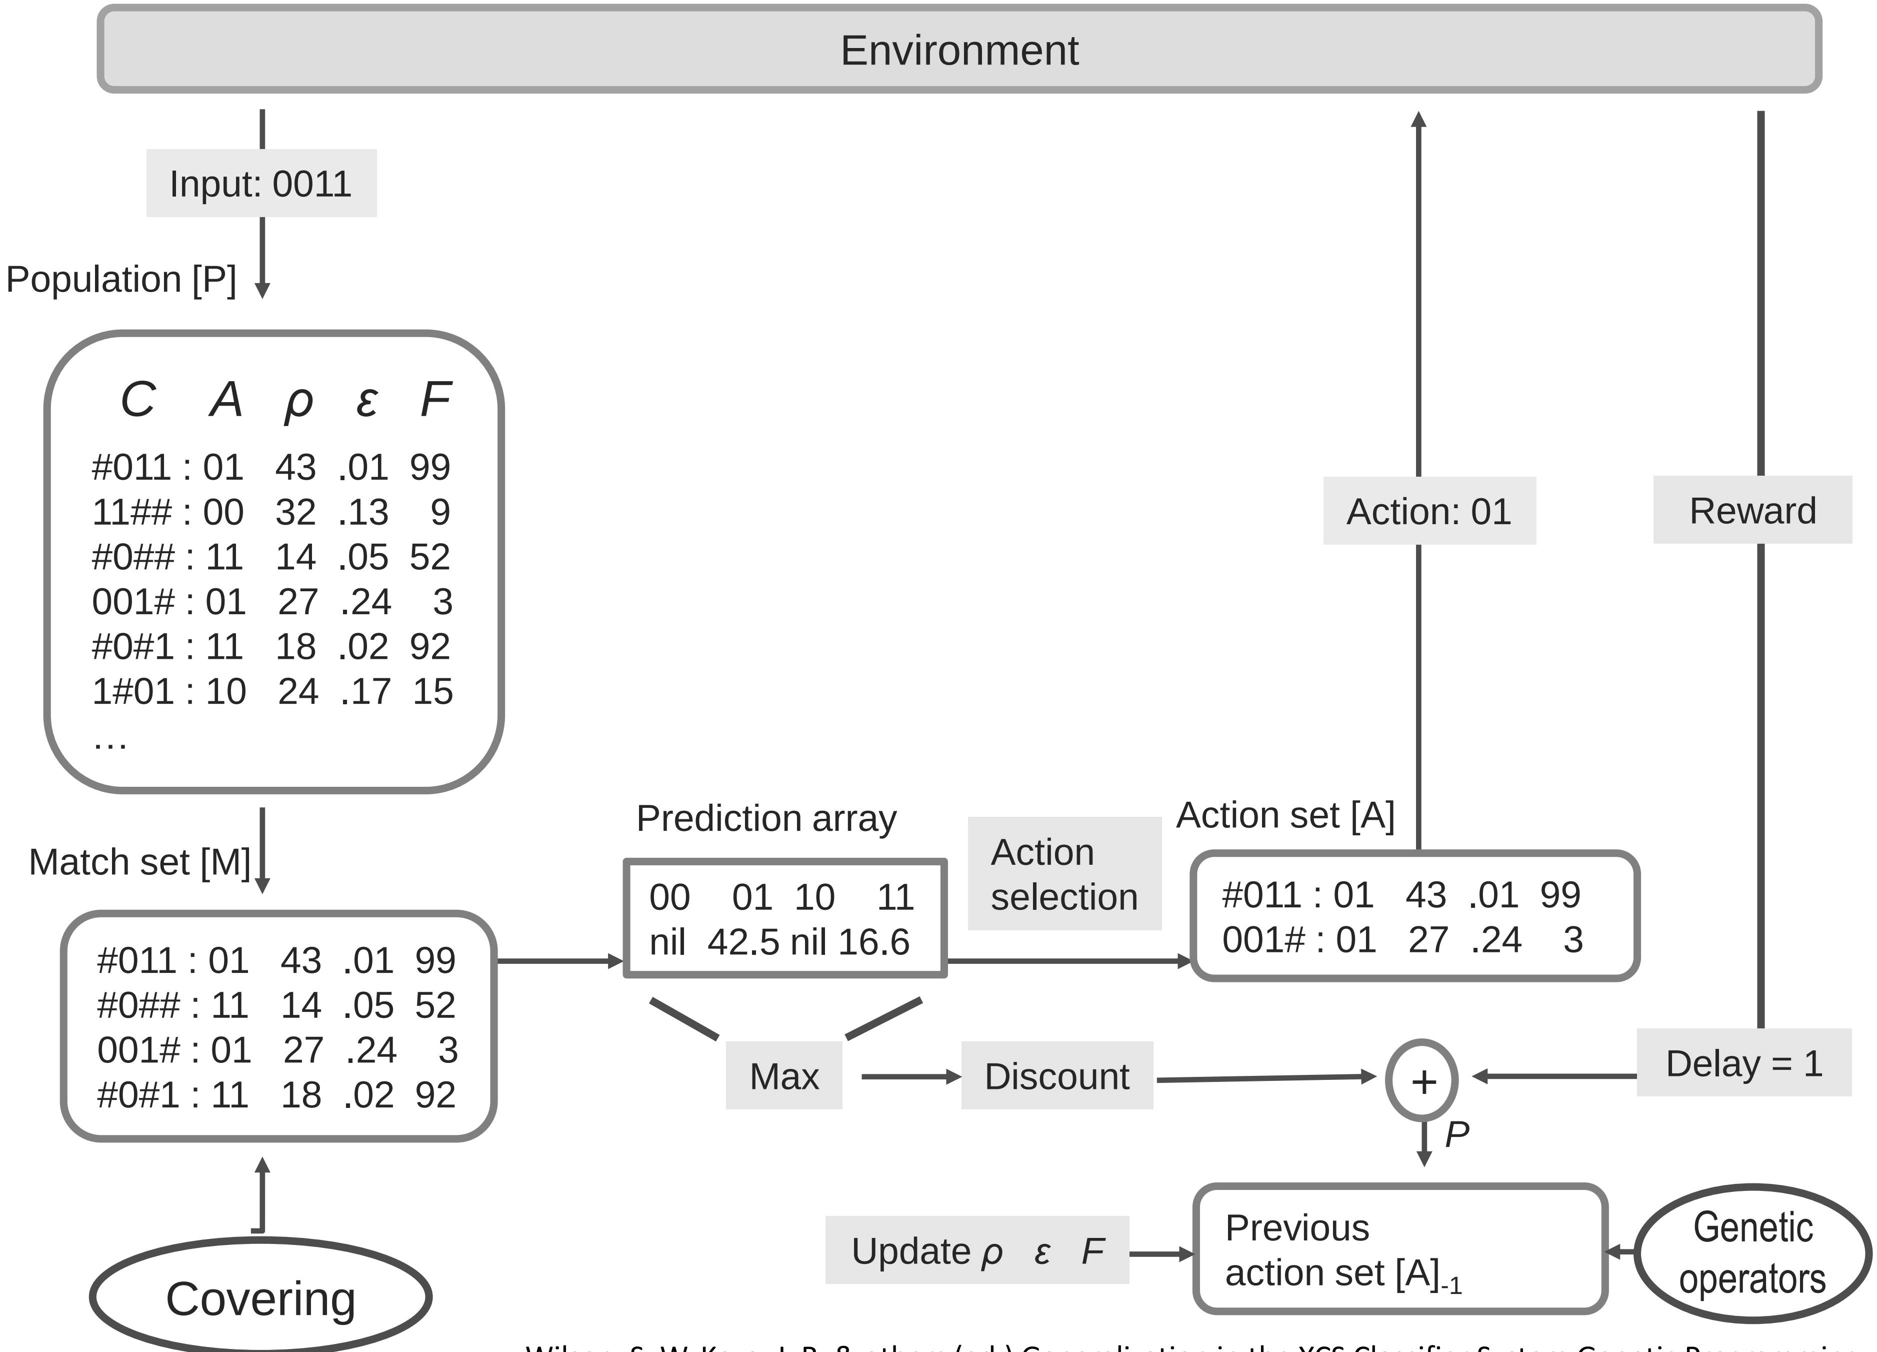

Wilson, S. W. Koza, J. R. & others (ed.) Generalization in the XCS Classifier System Genetic Programming 1998: Proceedings of the Third Annual Conference, 1998, 665-674

Learning Classifier Systems

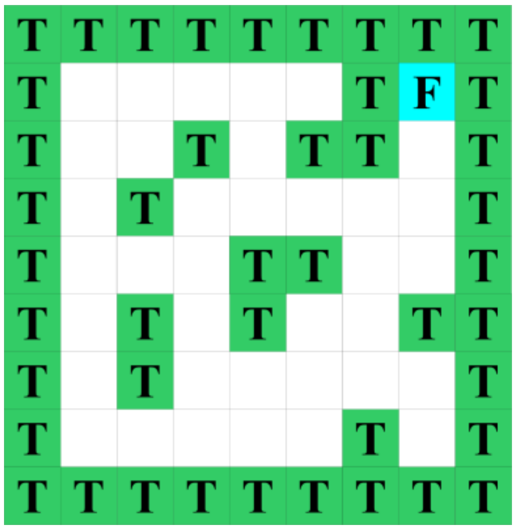

T: Tree
F: Food

- Standard (Multi-Step) Problem:
  - Maze6
- Goal:
  - Find shortest path to the food starting from a random point

Learning Classifier Systems

T	T	T	T	T	T	T	T	T
T						T	F	T
T			T		T	T		T
T		T						T
T				T	T			T
T		T		T			T	T
T		T						T
T						T		T
T	T	T	T	T	T	T	T	T

T: Tree
F: Food

- Problem:
  - Limited sensors, no global knowledge
    - Partially observable Markov decision process
- Solution:
  - Iterations, back-propagation of reward

Learning Classifier Systems

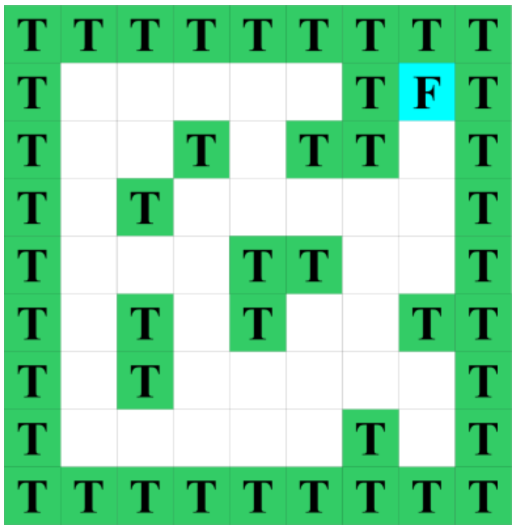

T: Tree
F: Food

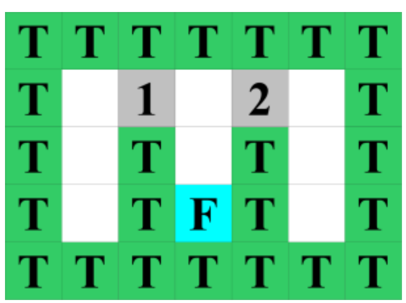

- Problem:
  - Limited sensors, no global knowledge
    - Partially observable Markov decision process
- Solution:
  - Iterations, back-propagation of reward
- Aliasing positions:
  - Handle by using memory

Learning Classifier Systems

- Many aliasing positions
- Other agents present
- Dynamic world
  - food and other agents move
- Limited sensors

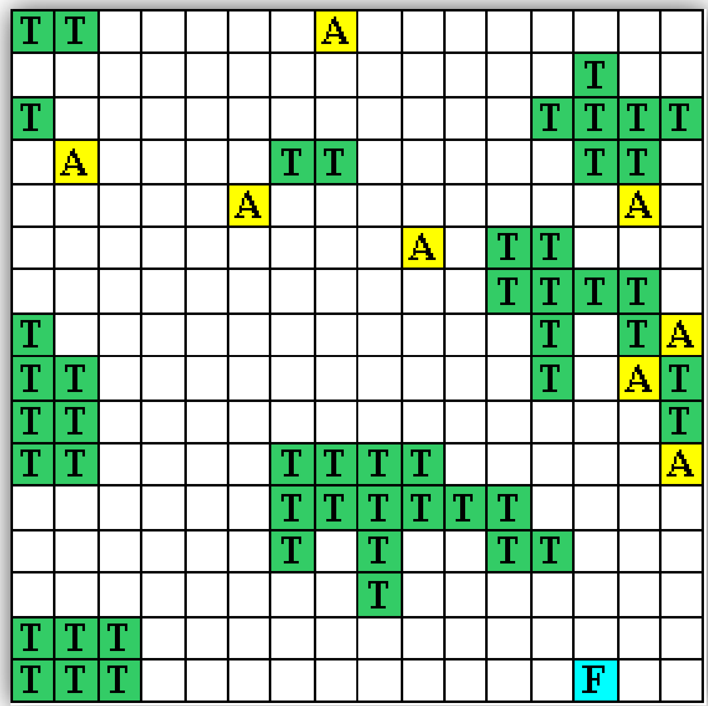

T: Trees
F: Food
A: Agent

Predator/Prey Scenarios

- Terminology:
  - Obstacles, Prey, Predator
- Goal: Get near the moving prey as often as possible
  - Global observation task
  - Continuous
  - Average Quality

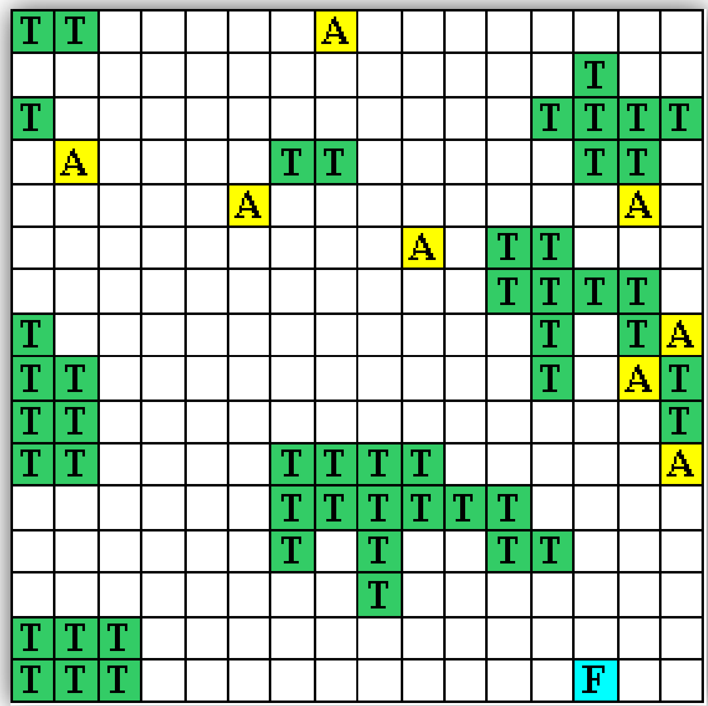

T: Trees/Obstacles
F: Food/Prey
A: Agent/Predator

Classification of Predator/Prey Scenarios

(1) Access to local information only
(2) open areas with some obstacles
(3) internal state unknown to others

- No standard MDP
  - Limited sensors (1, 3)
  - Aliasing positions (2)

Classification of Predator/Prey Scenarios

(1) Access to local information only
(2) open areas with some obstacles
(3) internal state unknown to others
(4) dynamic scenario

- No standard MDP
  - Limited sensors (1, 3)
  - Aliasing positions (2)
- No POMDP
  - Non-static scenario (4)

Classification of Predator/Prey Scenarios

(1) Access to local information only
(2) open areas with some obstacles
(3) internal state unknown to others
(4) dynamic scenario
(5) predators share global observation task
(6) runs continuously

- No standard MDP
  - Limited sensors (1, 3)
  - Aliasing positions (2)
- No POMDP
  - Non-static scenario (4)
- XCS has to be adapted
  - No "final" reward (5), no iterations (6)

Sensors

One sensor array for each direction

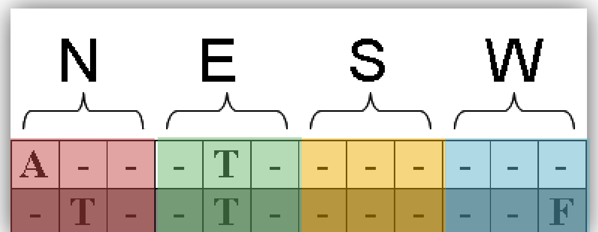

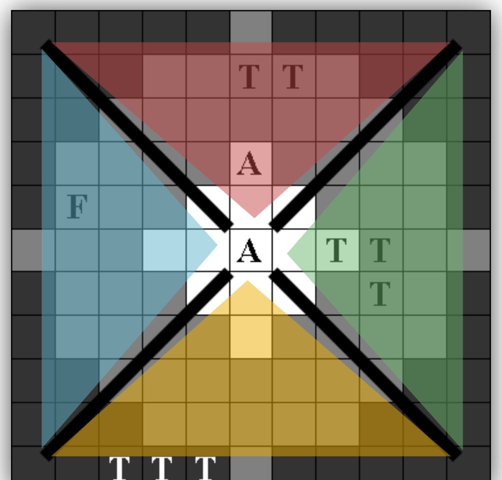

Sensors

One sensor array for each direction

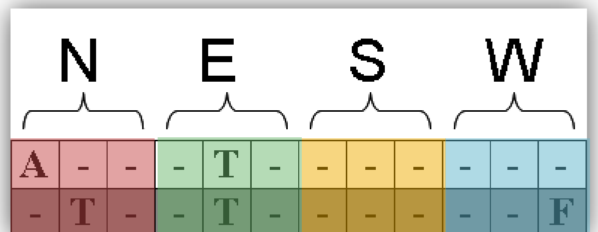

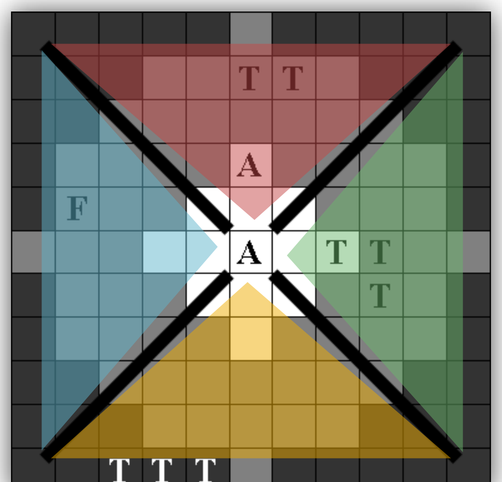

Sensors can sense either far or near

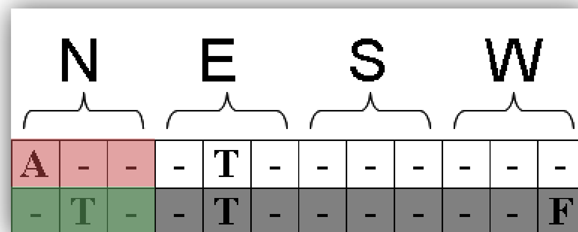

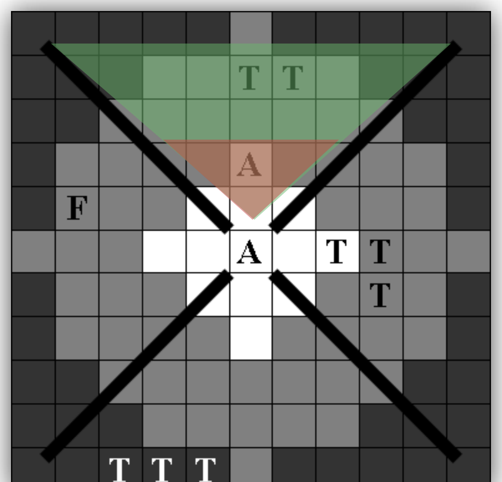

Sensors

One sensor array for each direction

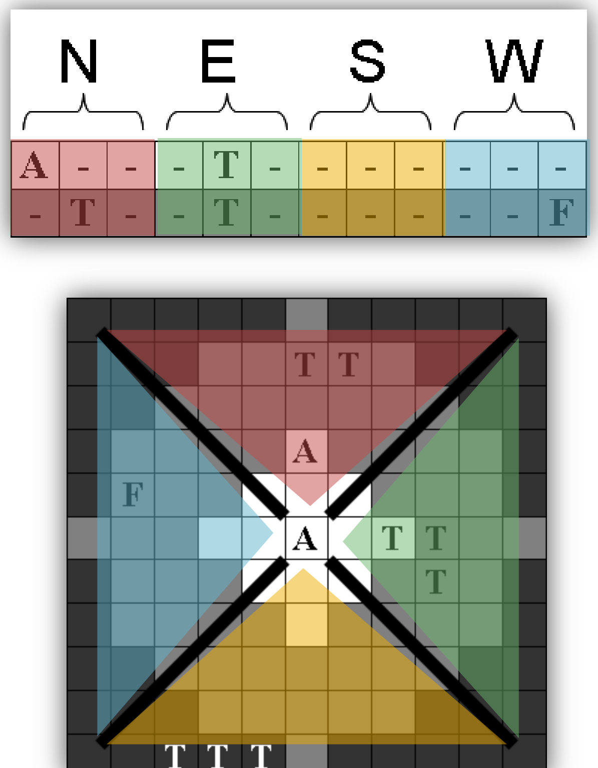

Sensors can sense either far or near (observation range / sight range)

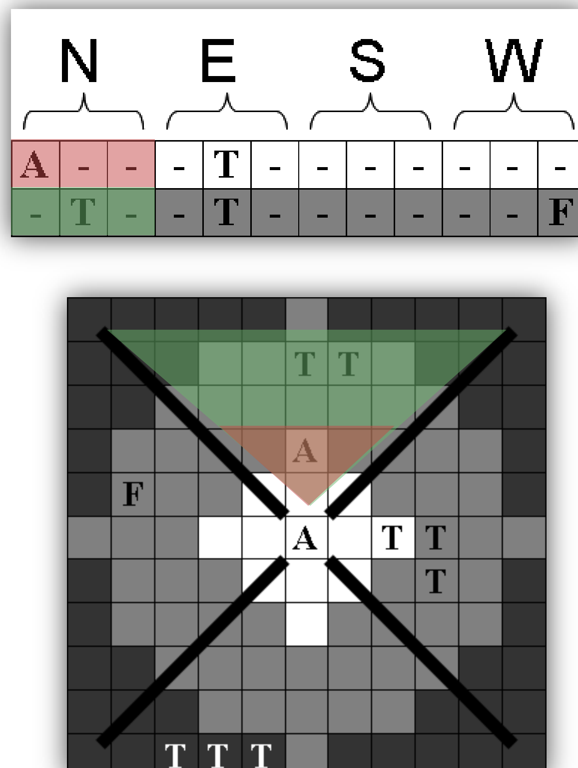

Sensors can distinguish between Predators, Prey and Obstacles

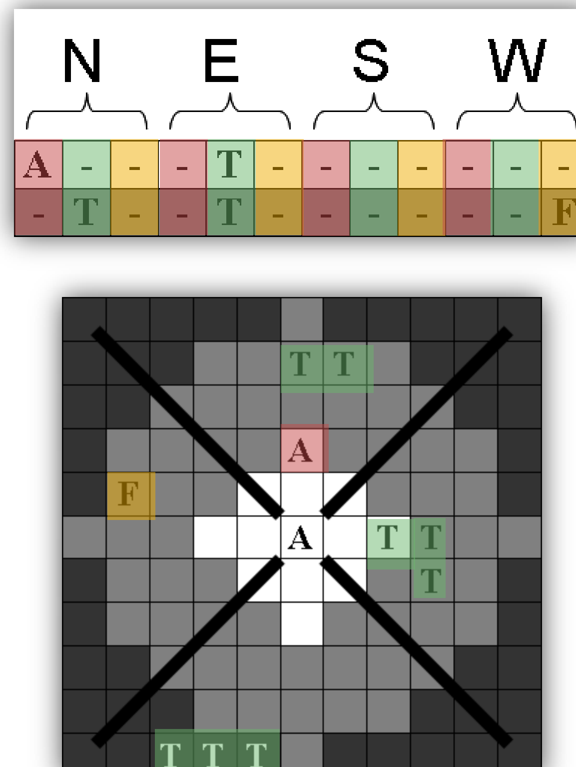

Adaption of the standard XCS Reward Function

- Standard implementation:

- Adapted implementation

Adaption of the standard XCS Reward Function

- Standard implementation:
  - Reward:
    - Food is in a neighboring cell
- Adapted implementation
  - Reward:
    - Food is in observation range ("XCS obs")
    - Food is in sight range ("XCS sight")

Adaption of the standard XCS Reward Function

- Standard implementation:
  - Reward:
    - Food is in a neighboring cell
  - Action:
    - Assign reward
    - Restart scenario
    - Switch between explore/exploit phase
- Adapted implementation
  - Reward:
    - Food is in observation range ("XCS obs")
    - Food is in sight range ("XCS sight")
  - Action:
    - Assign reward
    - Continue scenario
    - Always use exploit phase

Classification of Predator/Prey Scenarios

- Global knowledge cannot be reconstructed
  - Memory becomes invalid after each step

- The predator/prey scenario is a Non-observable Markov Decision Processes

Classification of Predator/Prey Scenarios

- Global knowledge cannot be reconstructed
  - Memory becomes invalid after each step
- The predator/prey scenario is a Non-observable Markov Decision Processes
- Despite being a NOMDP, can the XCS still learn?

Classification of Predator/Prey Scenarios

- Global knowledge cannot be reconstructed
  - Memory becomes invalid after each step
- The predator/prey scenario is a Non-observable Markov Decision Processes
- Despite being a NOMDP, can the XCS still learn?
  - Yes, the prey does move only a limited number of cells

    → The environment does not completely change each turn

Testing Methodology

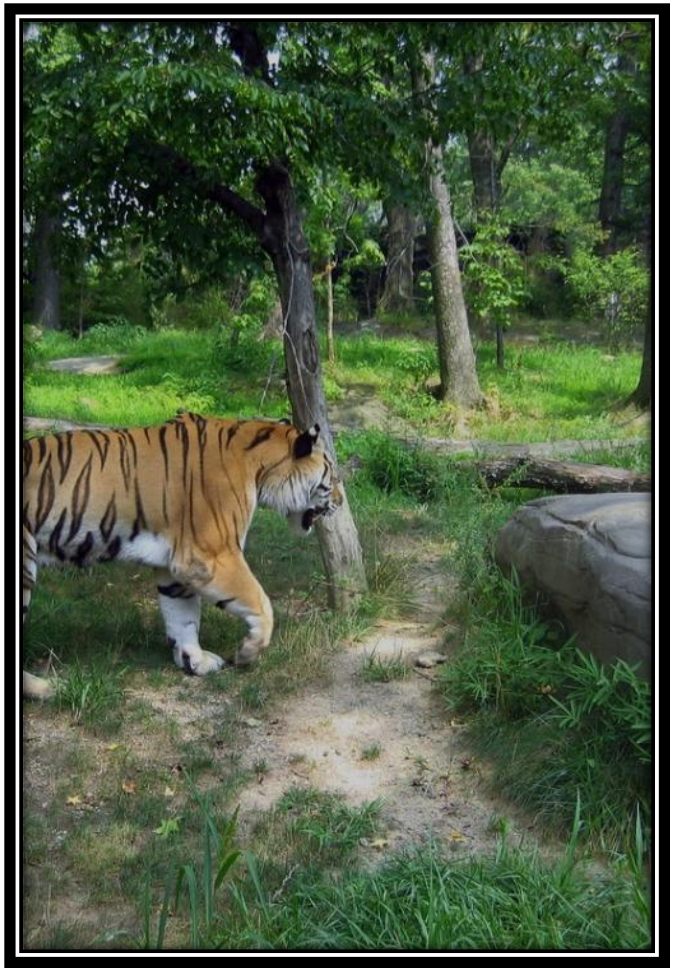

http://www.flickr.com/photos/james_crowley

- “XCS obs”, “XCS sight”
- “Obstacle-evading prey”
- “Predator-evading prey”
- “Blinded Prey”
- Standard XCS parameter settings

Scenario Configurations

- Three different scenarios with observation and sight ranges

"Pillar scenario"

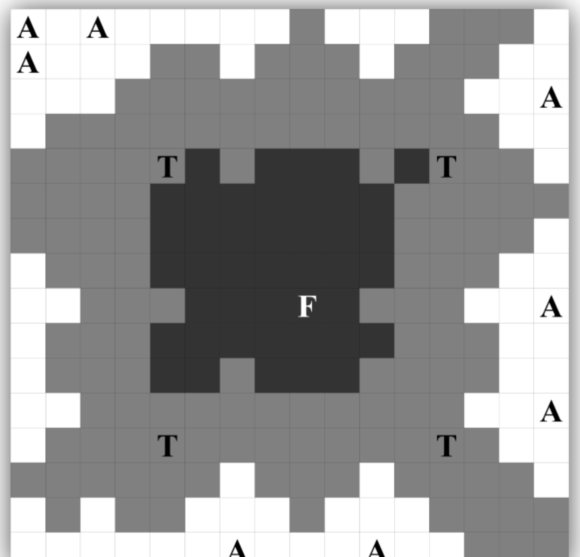

"Random scenario"

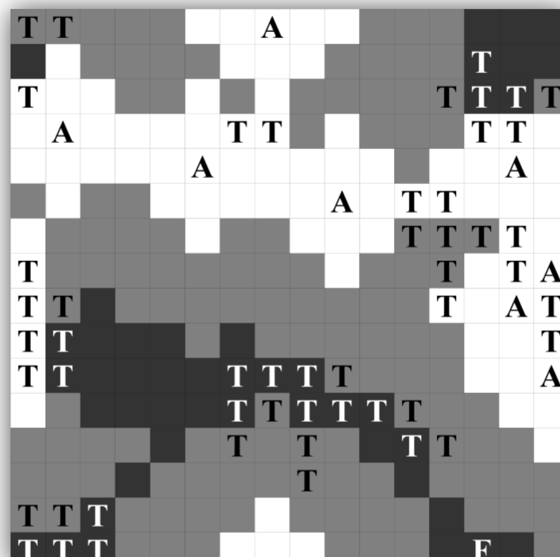

"Difficult scenario"

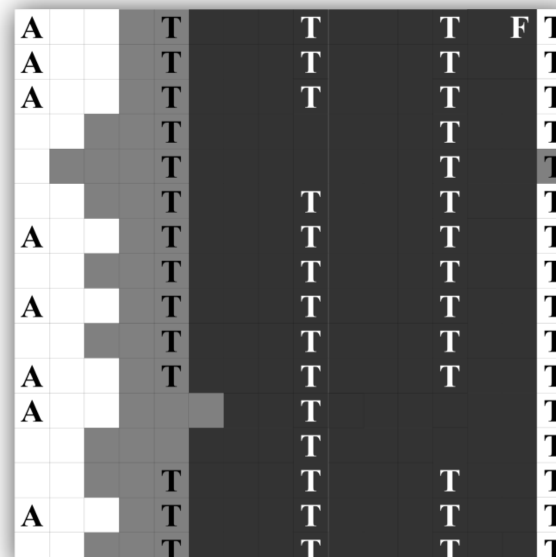

XCS Experimental Results "Pillar Scenario"

- XCS (obs) shows some learning
- XCS (sight)

Obstacle-evading prey

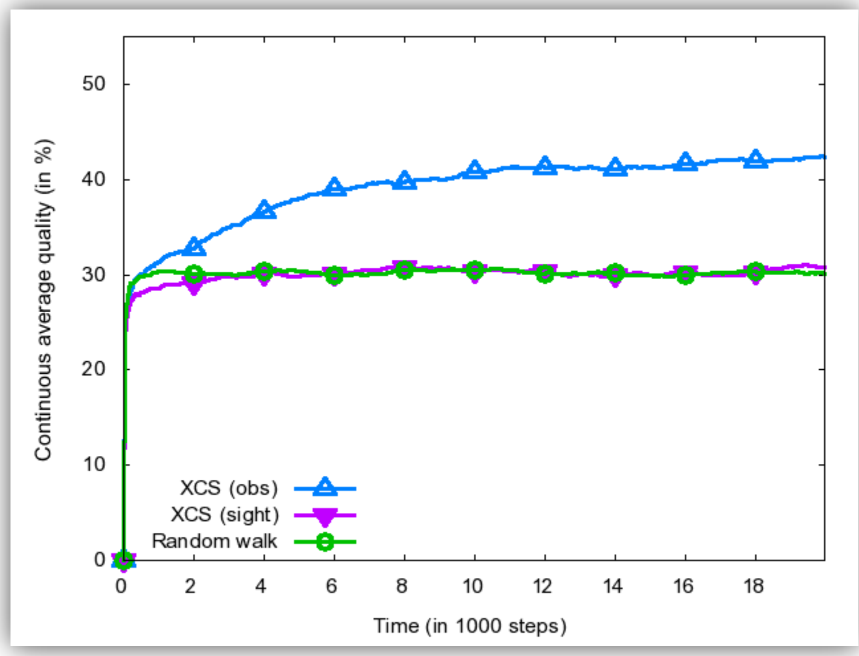

Predator-evading prey

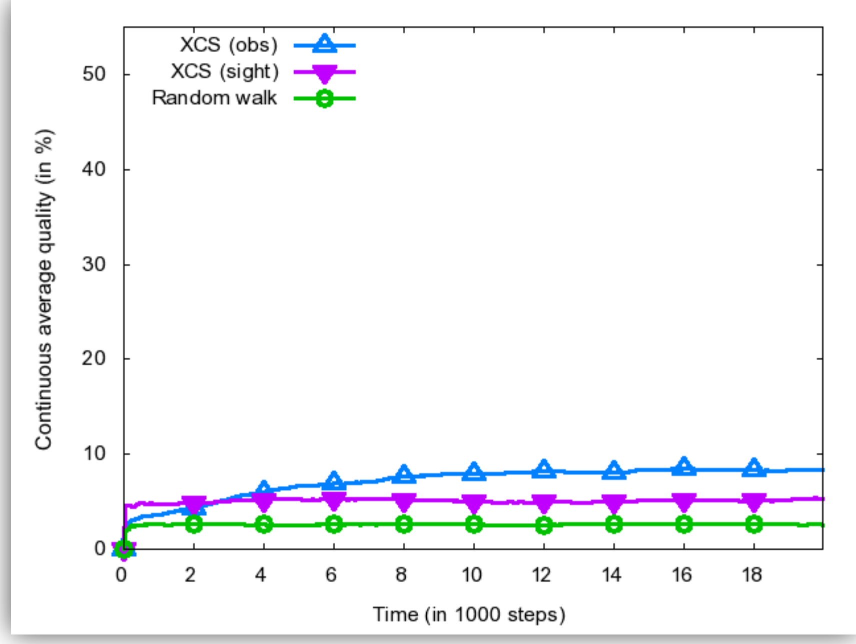

XCS Experimental Results "Random Scenario"

- XCS shows very little learning

Obstacle-evading prey

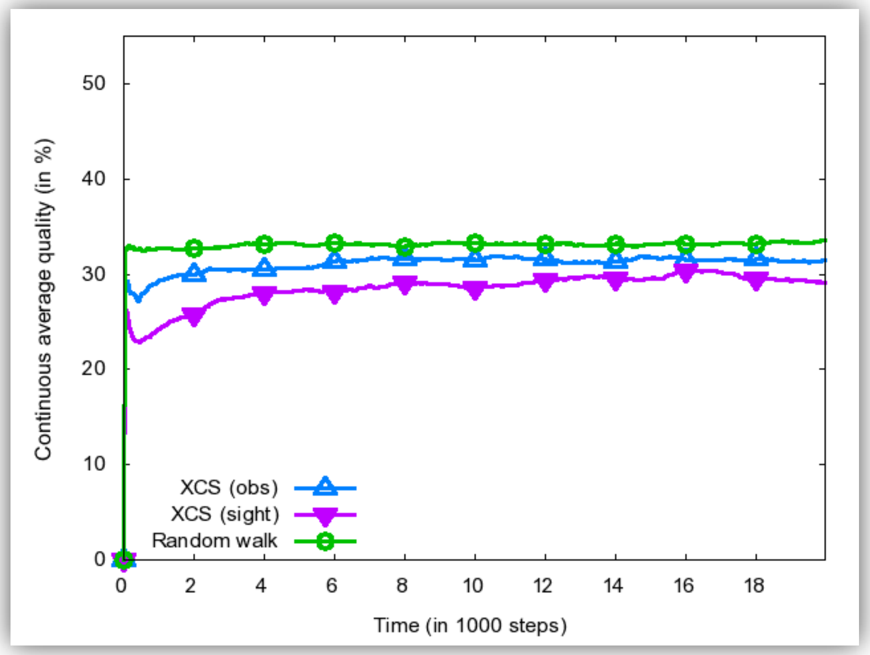

Predator-evading prey

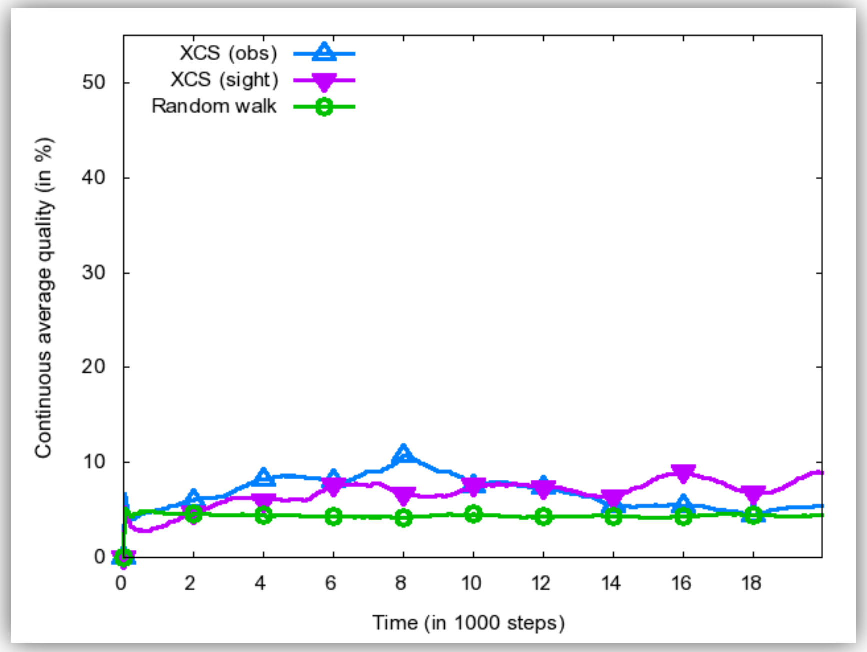

XCS Experimental Results "Difficult Scenario"

- XCS shows significant learning
  - But also unlearning after 8,000 steps
- "Difficult Scenario" is a maze-like scenario, this result was expected

Blinded prey

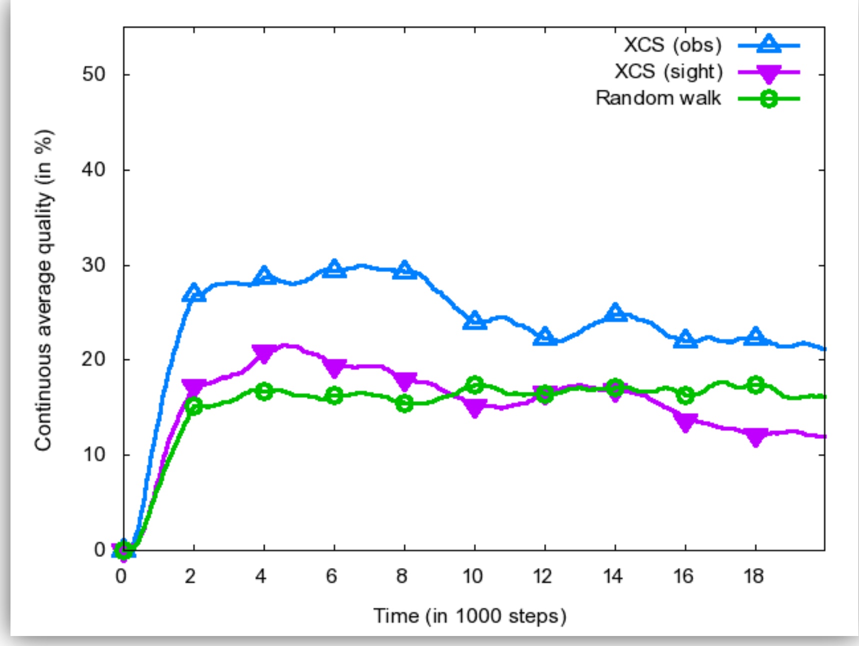

Reward Events “eventXCS”

	N	E	S	W
A	1	0	0	0
T	0	1	0	0
F	0	0	0	0
A	0	0	0	0
T	1	1	0	0
F	0	0	0	1

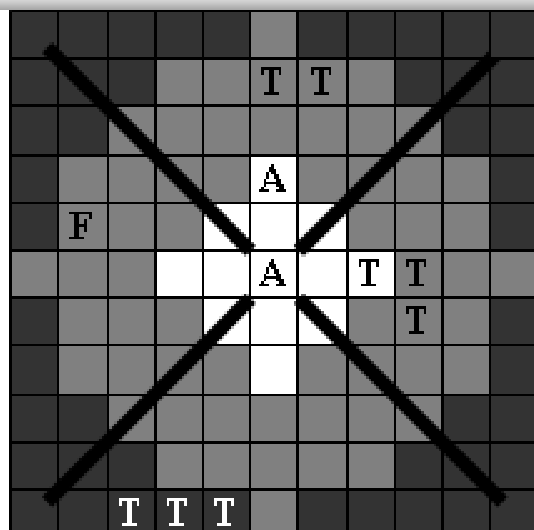

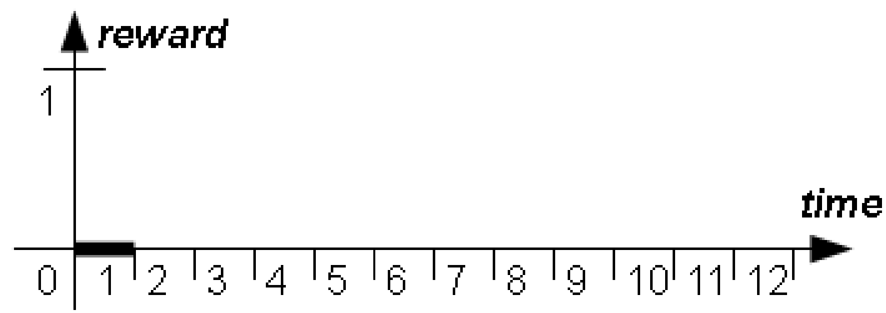

Reward Events "eventXCS"

	N	E	S	W
A	0	0	0	0
T	0	0	0	0
F	0	0	0	0
A	1	0	0	0
T	1	1	1	0
F	0	0	0	1

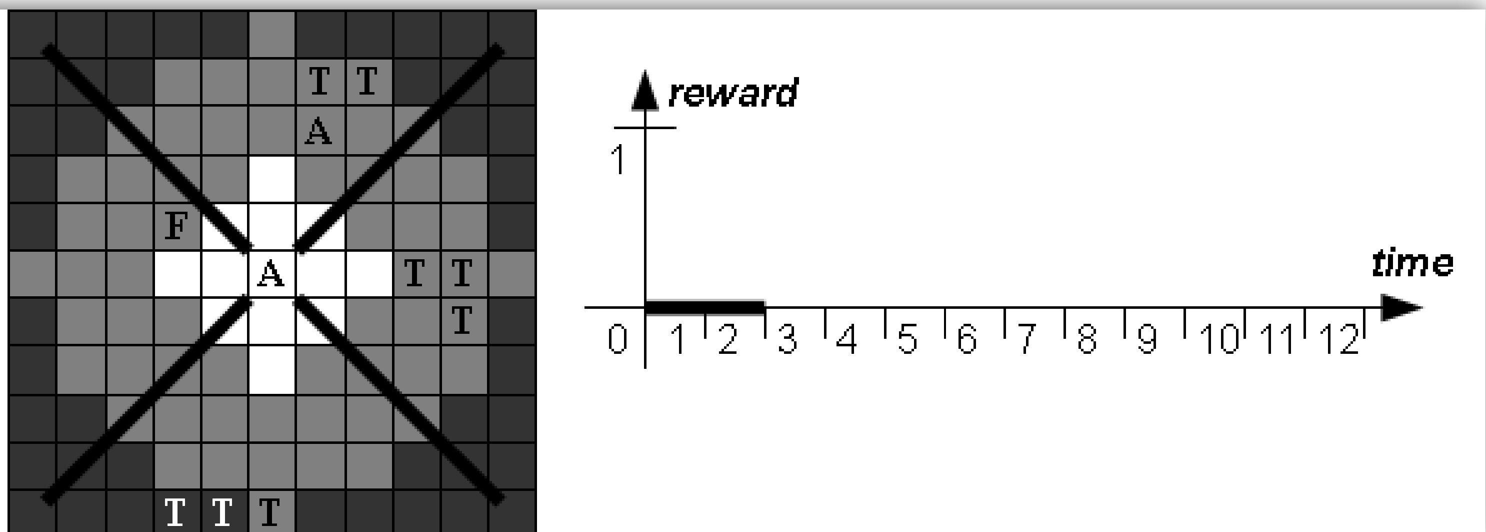

Reward Events “eventXCS”

	N	E	S	W
A	0	0	0	0
T	0	0	0	0
F	0	0	0	0
A	1	0	0	0
T	1	1	1	0
F	1	0	0	0

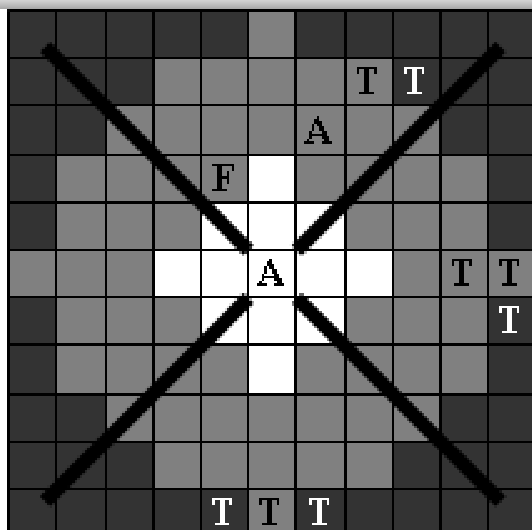

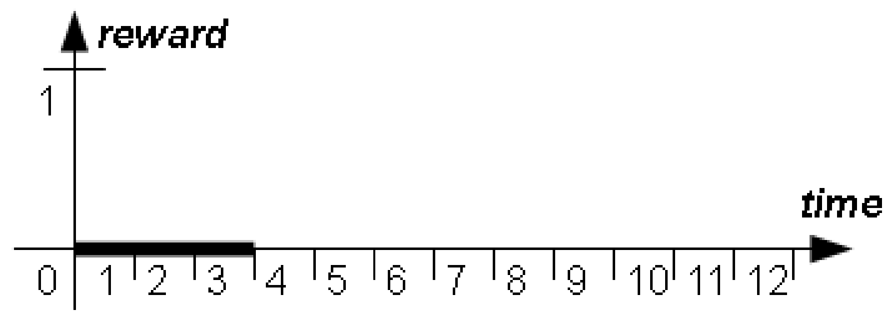

Reward Events “eventXCS”

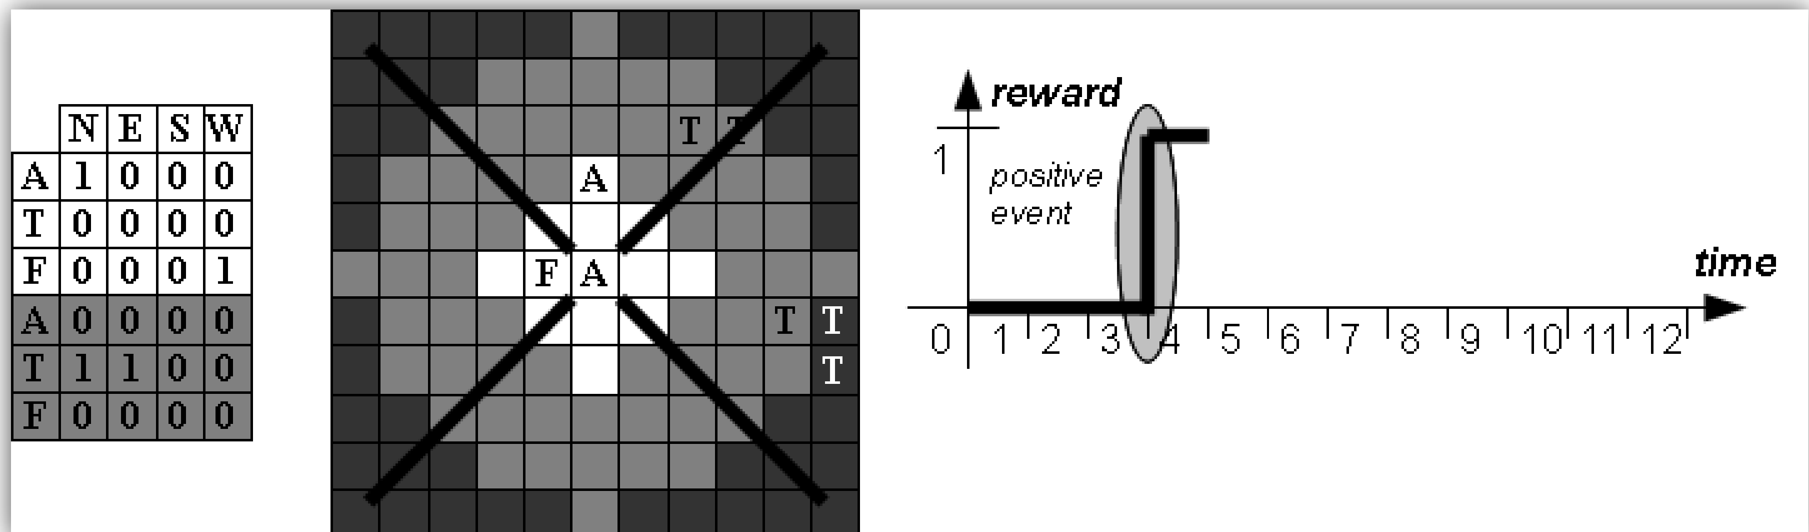

Reward Events "eventXCS"

	N	E	S	W
A	0	0	0	0
T	0	0	0	0
F	1	0	0	0
A	1	0	0	0
T	1	0	0	0
F	0	0	0	0

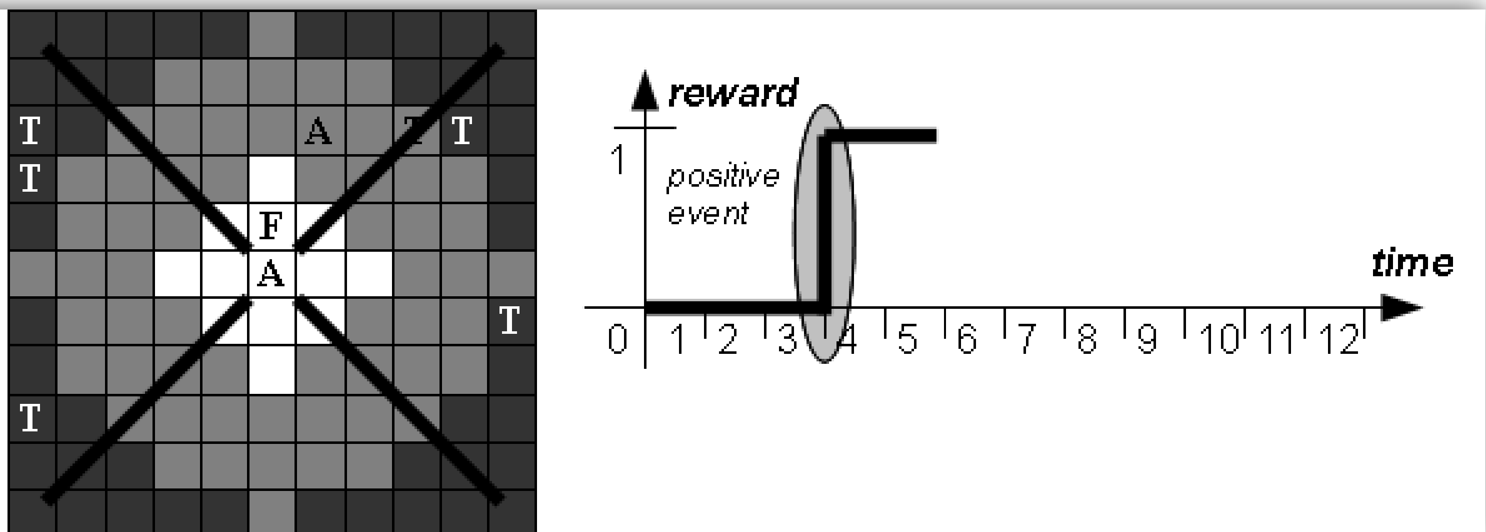

Reward Events "eventXCS"

	N	E	S	W
A	0	0	0	0
T	0	0	0	0
F	1	0	0	0
A	1	0	0	0
T	0	1	0	0
F	0	0	0	0

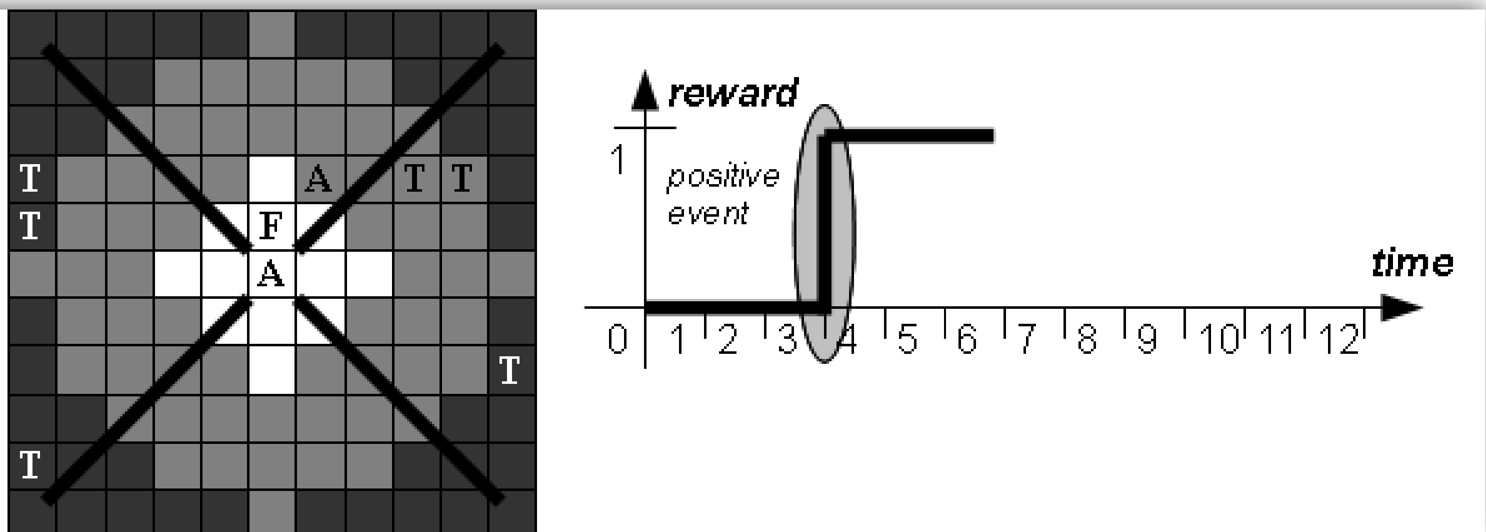

Reward Events "eventXCS"

	N	E	S	W
A	0	0	0	0
T	0	0	0	0
F	1	0	0	0
A	1	0	0	0
T	0	1	0	1
F	0	0	0	0

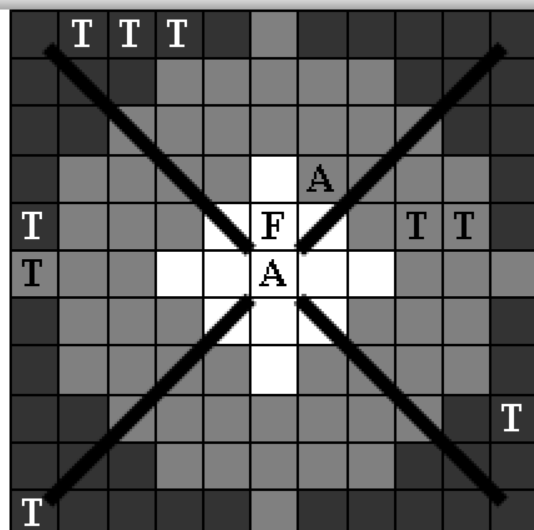

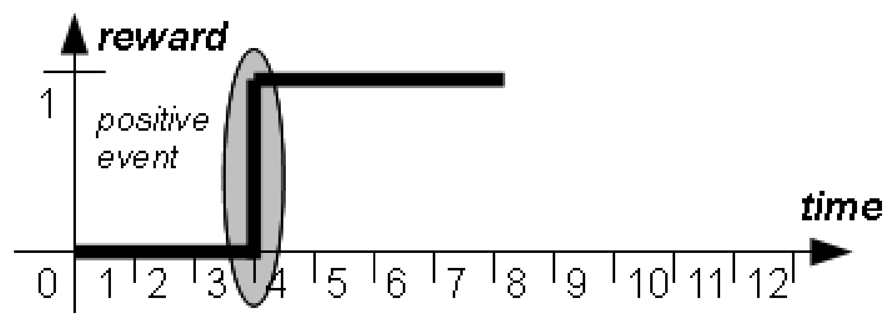

Reward Events “eventXCS”

	N	E	S	W
A	0	0	0	0
T	0	0	0	0
F	0	0	0	0
A	1	0	0	0
T	0	1	0	1
F	0	1	0	0

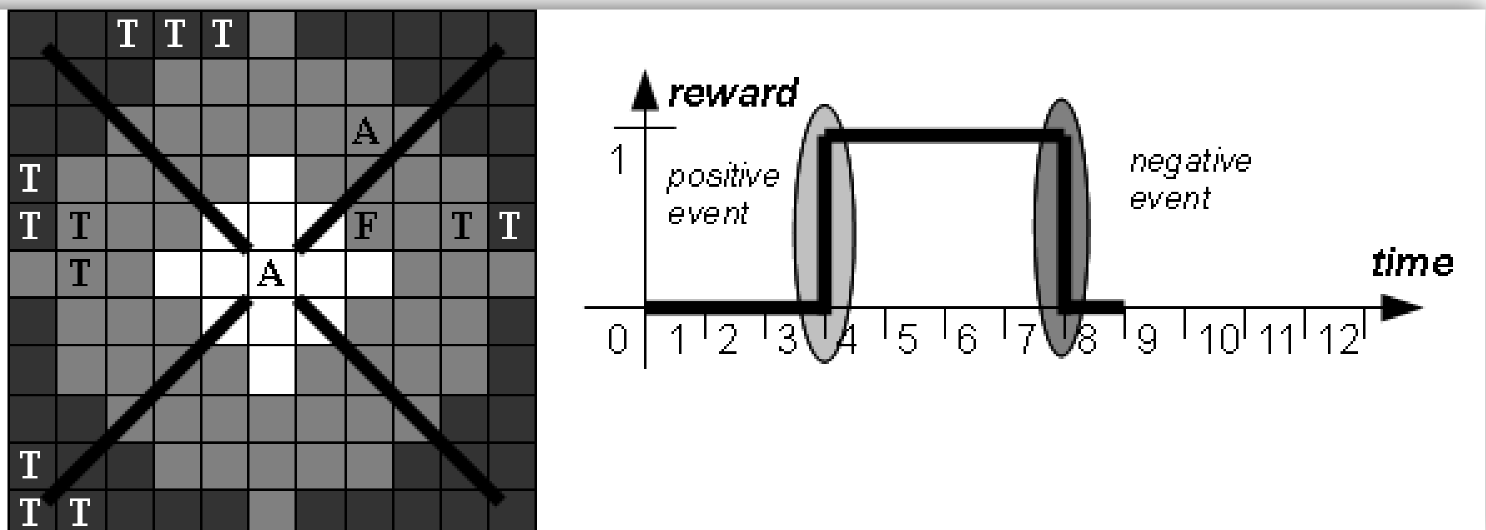

Reward Distribution "eventXCS"

- Analyze succession of positive and negative events

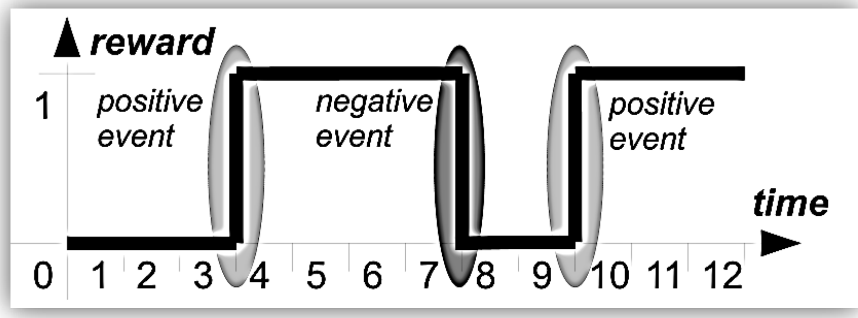

Reward Distribution "eventXCS"

- Analyze succession of positive and negative events
- Distribute the reward as soon as possible (i.e. at each event)
- Idea:
  - Action sets closer to an event probably contributed more

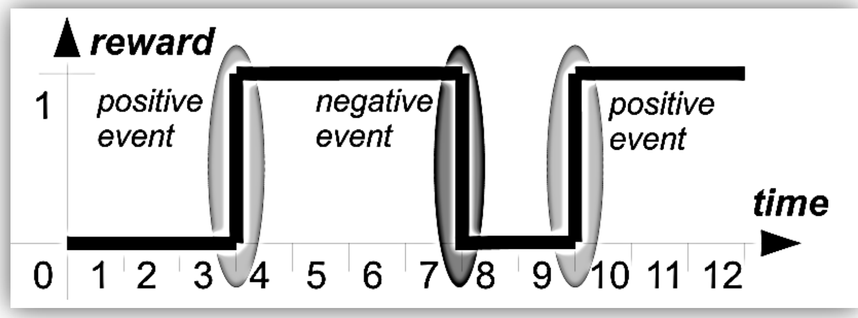

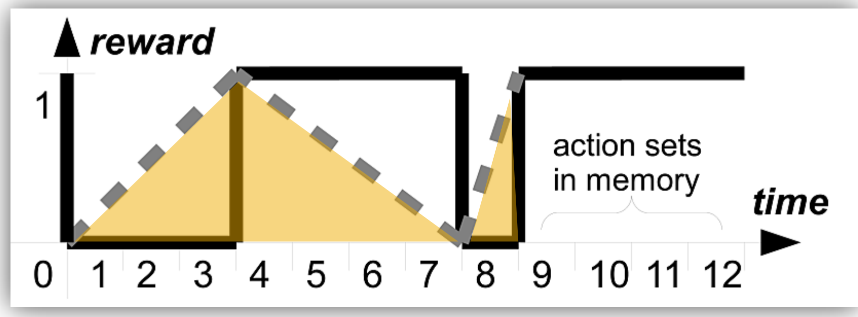

eventXCS Experimental Results "Pillar Scenario"

- eventXCS clearly outperforms XCS

Obstacle-evading prey

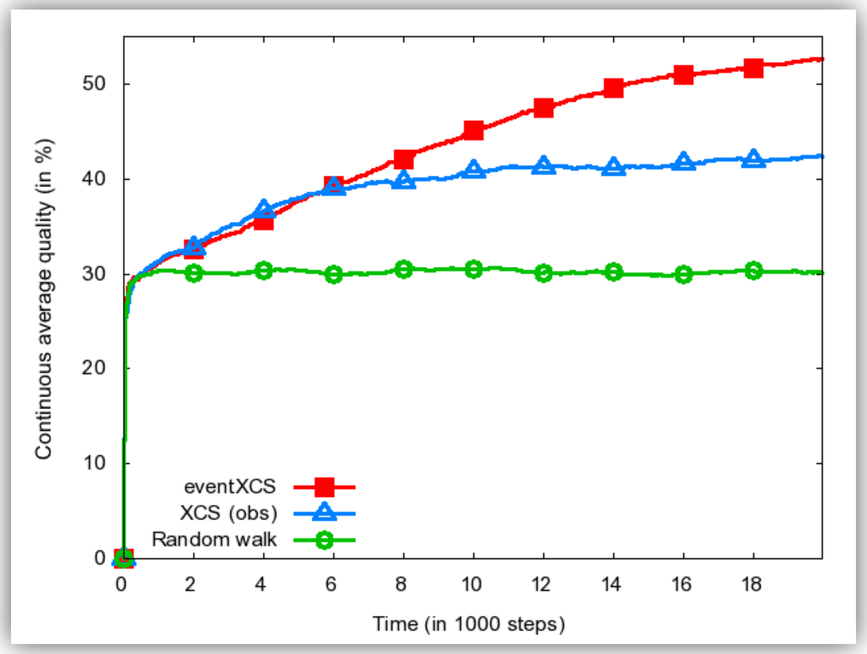

Predator-evading prey

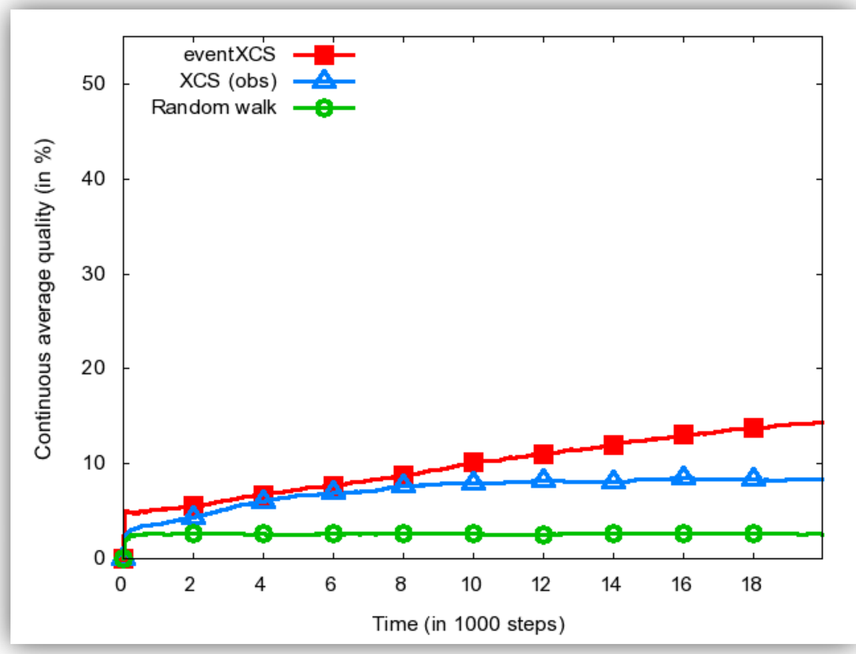

eventXCS Experimental Results "Random Scenario"

- eventXCS shows slow but steady learning with an obstacle-evading prey

Obstacle-evading prey

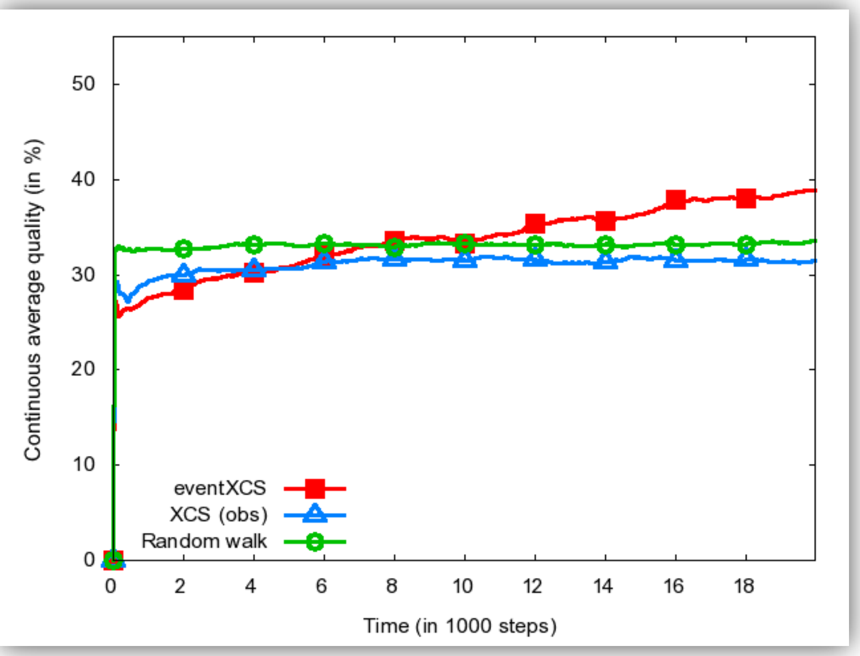

Predator-evading prey

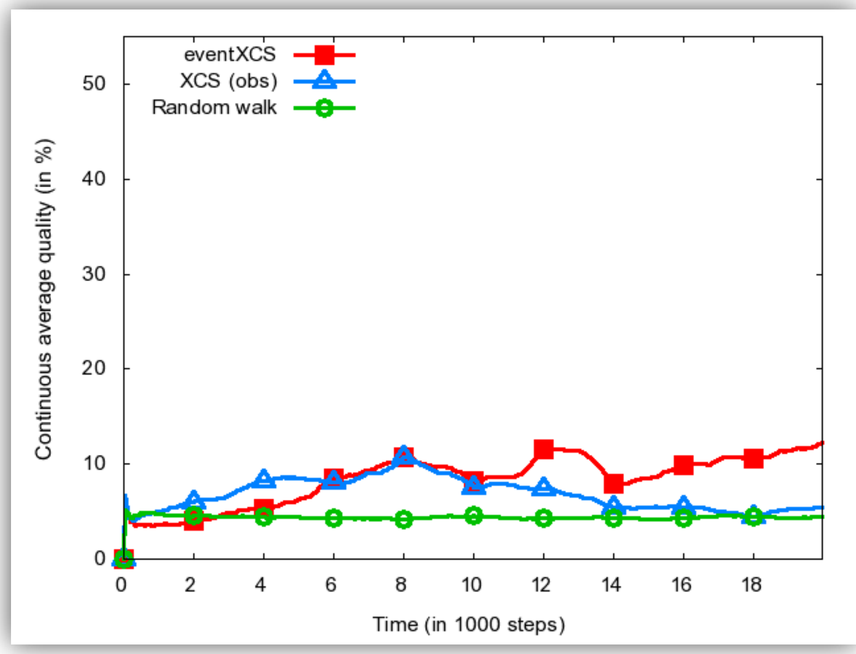

eventXCS Experimental Results "Difficult Scenario"

- eventXCS fails in this scenario
- Using "tournament selection" shows acceptable results with no sign of unlearning

Blinded prey

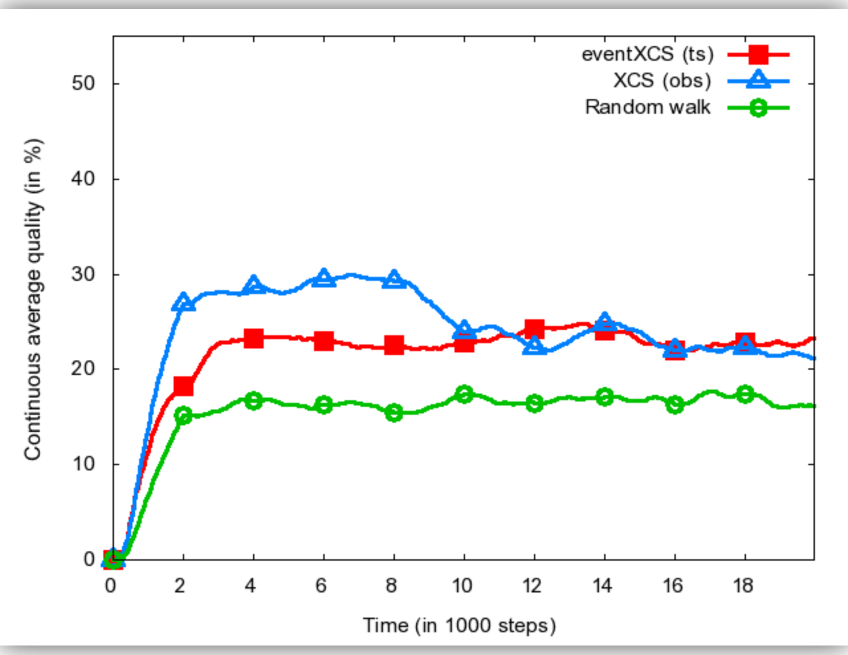

Conclusion

- XCS with minimal adaptions can learn
  - Unable to use sight range
- Event XCS superior

Adaption of XCS to Multi-Learner Predator/Prey Scenarios

- Thank you for your attention!

Backup slides

Neutral Events

- Neutral Event
  - No positive or negative event for a number of steps
  - Half of the action sets is discarded and receives reward
  - Idea:
    - Good actions are rewarded earlier
    - Preventing of dead ends

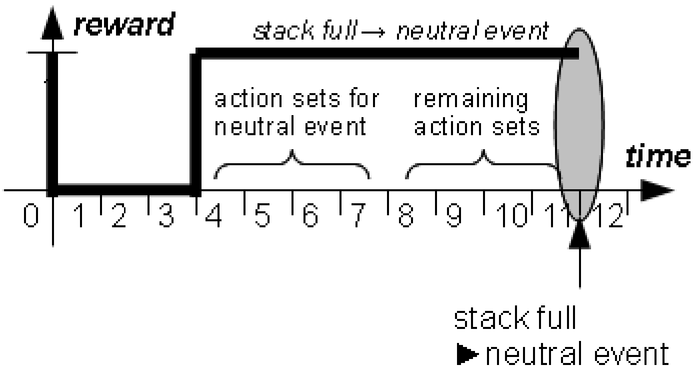

Neutral Events

- Neutral Event
  - No positive or negative event for a number of steps
  - Half of the action sets is discarded and receives reward
  - Idea:
    - Good actions are rewarded earlier
    - Preventing of dead ends
  - Problem:
    - Error possibility high if directly followed by an event.

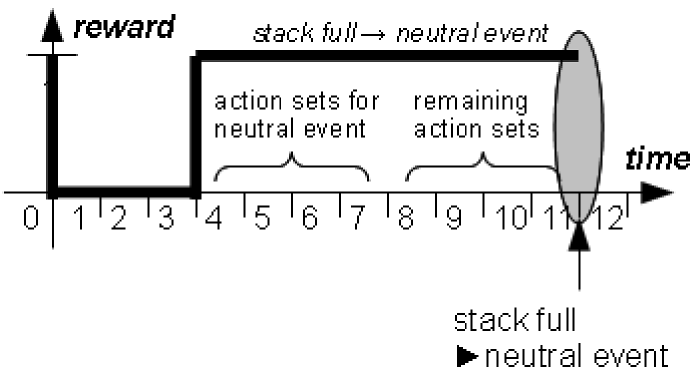

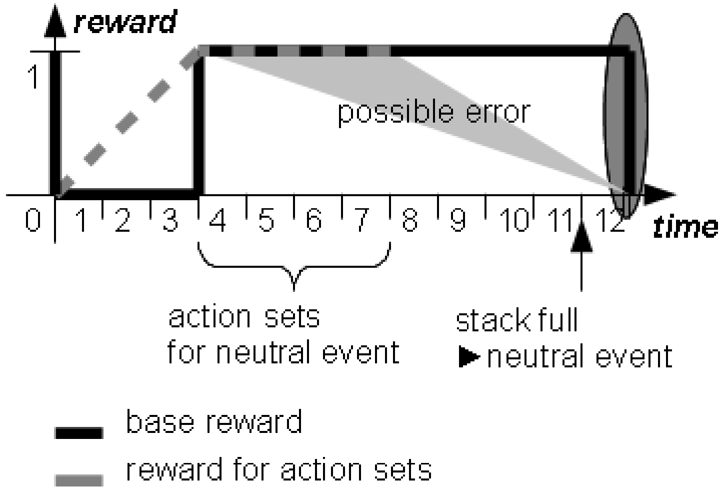

Neutral Events

- Tests have shown that a stack size of 8 is generally good for all three scenarios

"Pillar Scenario"

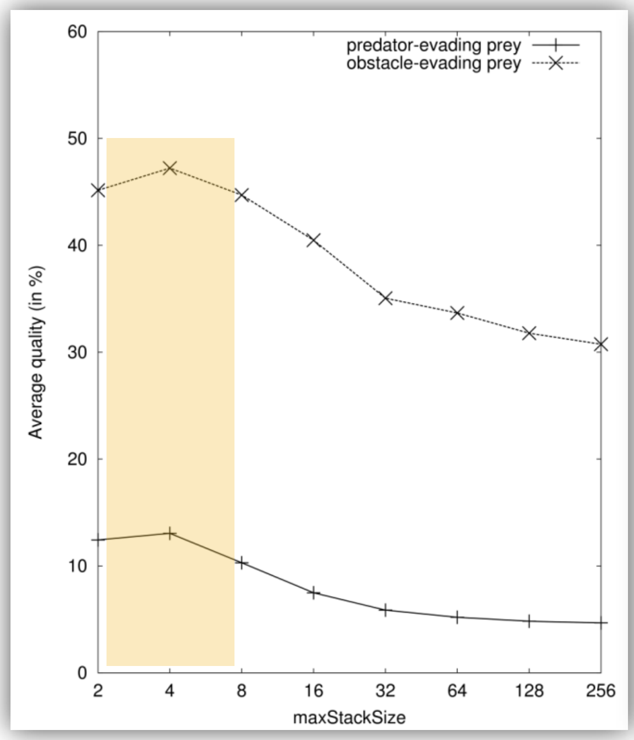

"Random Scenario"

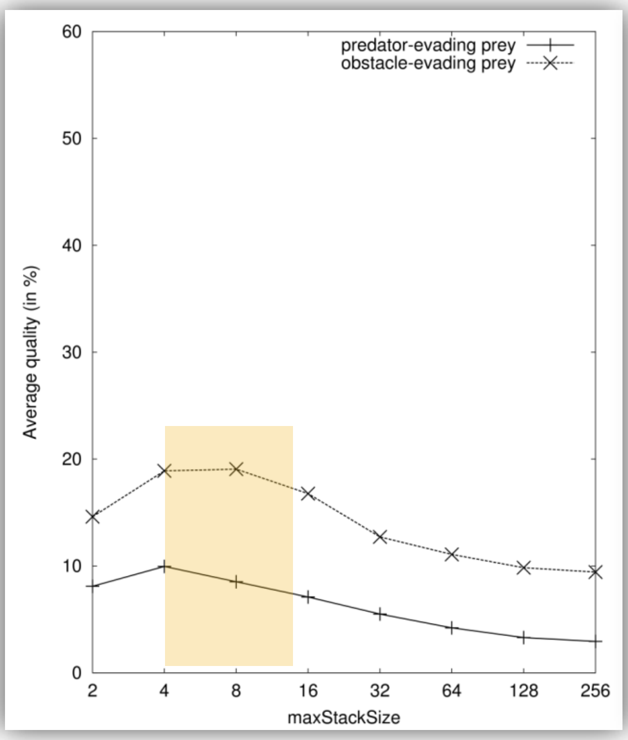

"Difficult Scenario"

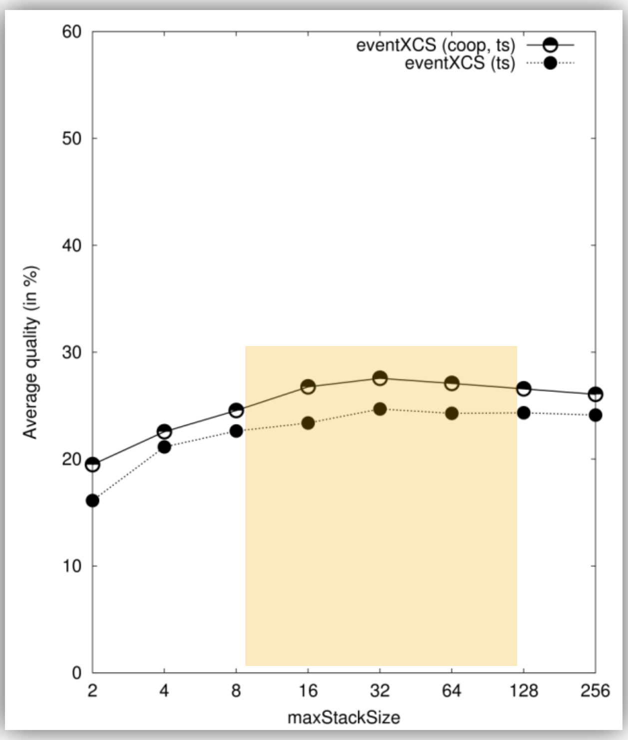

Learning Rate β

- Pillar Scenario
  - Obstacle-evading prey
  - Low learning rate (0.05) good, eventXCS very stable

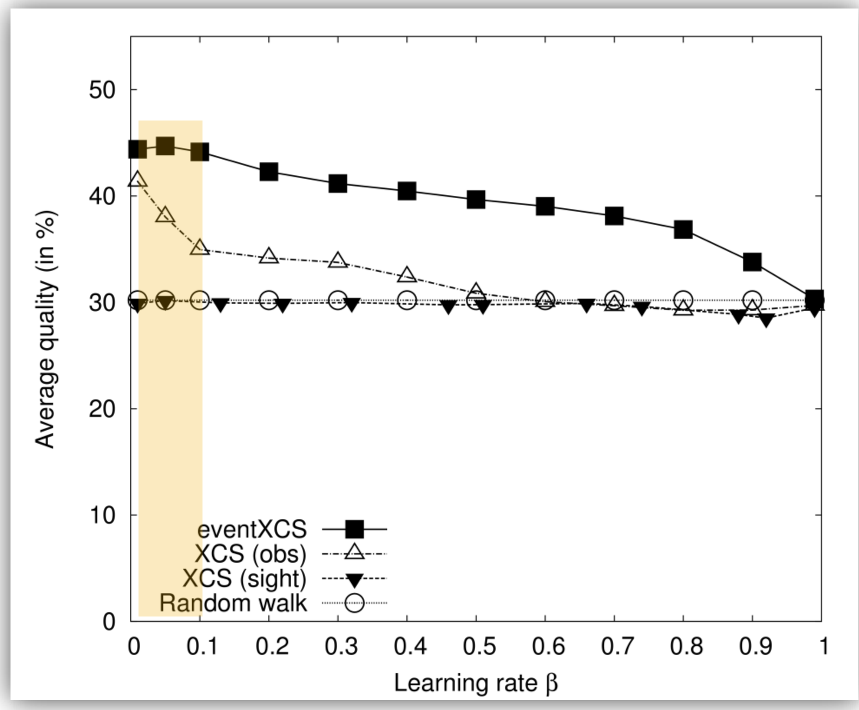

Learning Rate β

Pillar Scenario
Predator-evading
prey

Random Scenario,
Obstacle evading prey

Random Scenario,
Predator evading

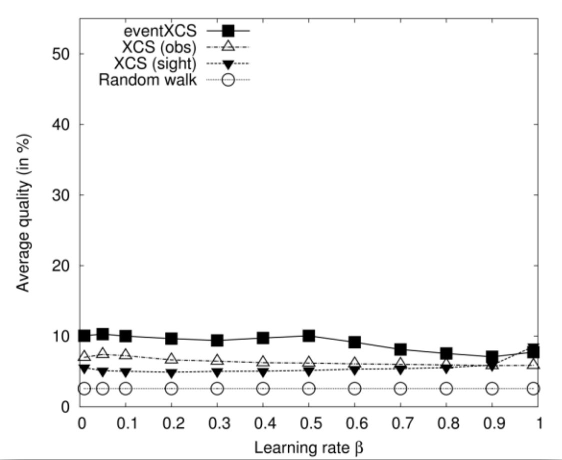

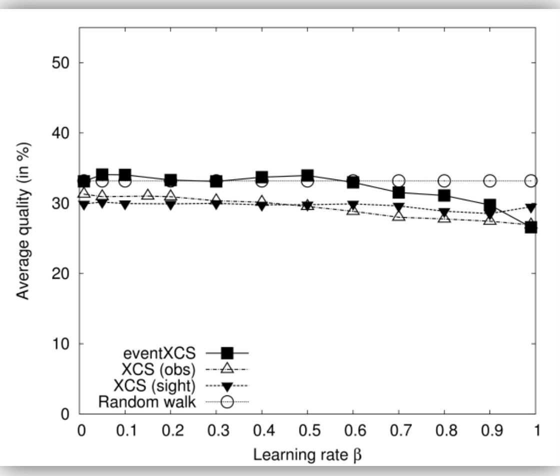

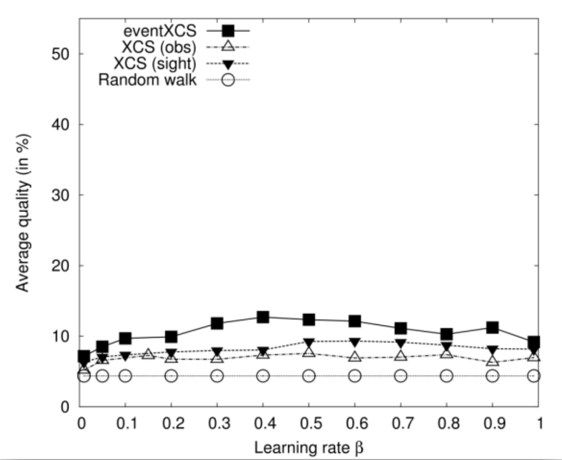

Learning Rate β

- Difficult Scenario
  - Blind prey

  - High learning rates show an advantage because of long distance to the prey