Clemens Lode, Urban Richter, Hartmut Schmeck
Karlsruhe Institute of Technology (Germany)
Institute AIFB
July, GECCO 2010

Adaption of XCS to Multi-Learner Predator/Prey Scenarios

http://www.flickr.com/photos/shreeram

Outline

http://www.flickr.com/photos/yathin

- Learning Classifier Systems
- XCS in Predator/Prey Scenarios
- Adapting the Reward Function
- Experimental Results

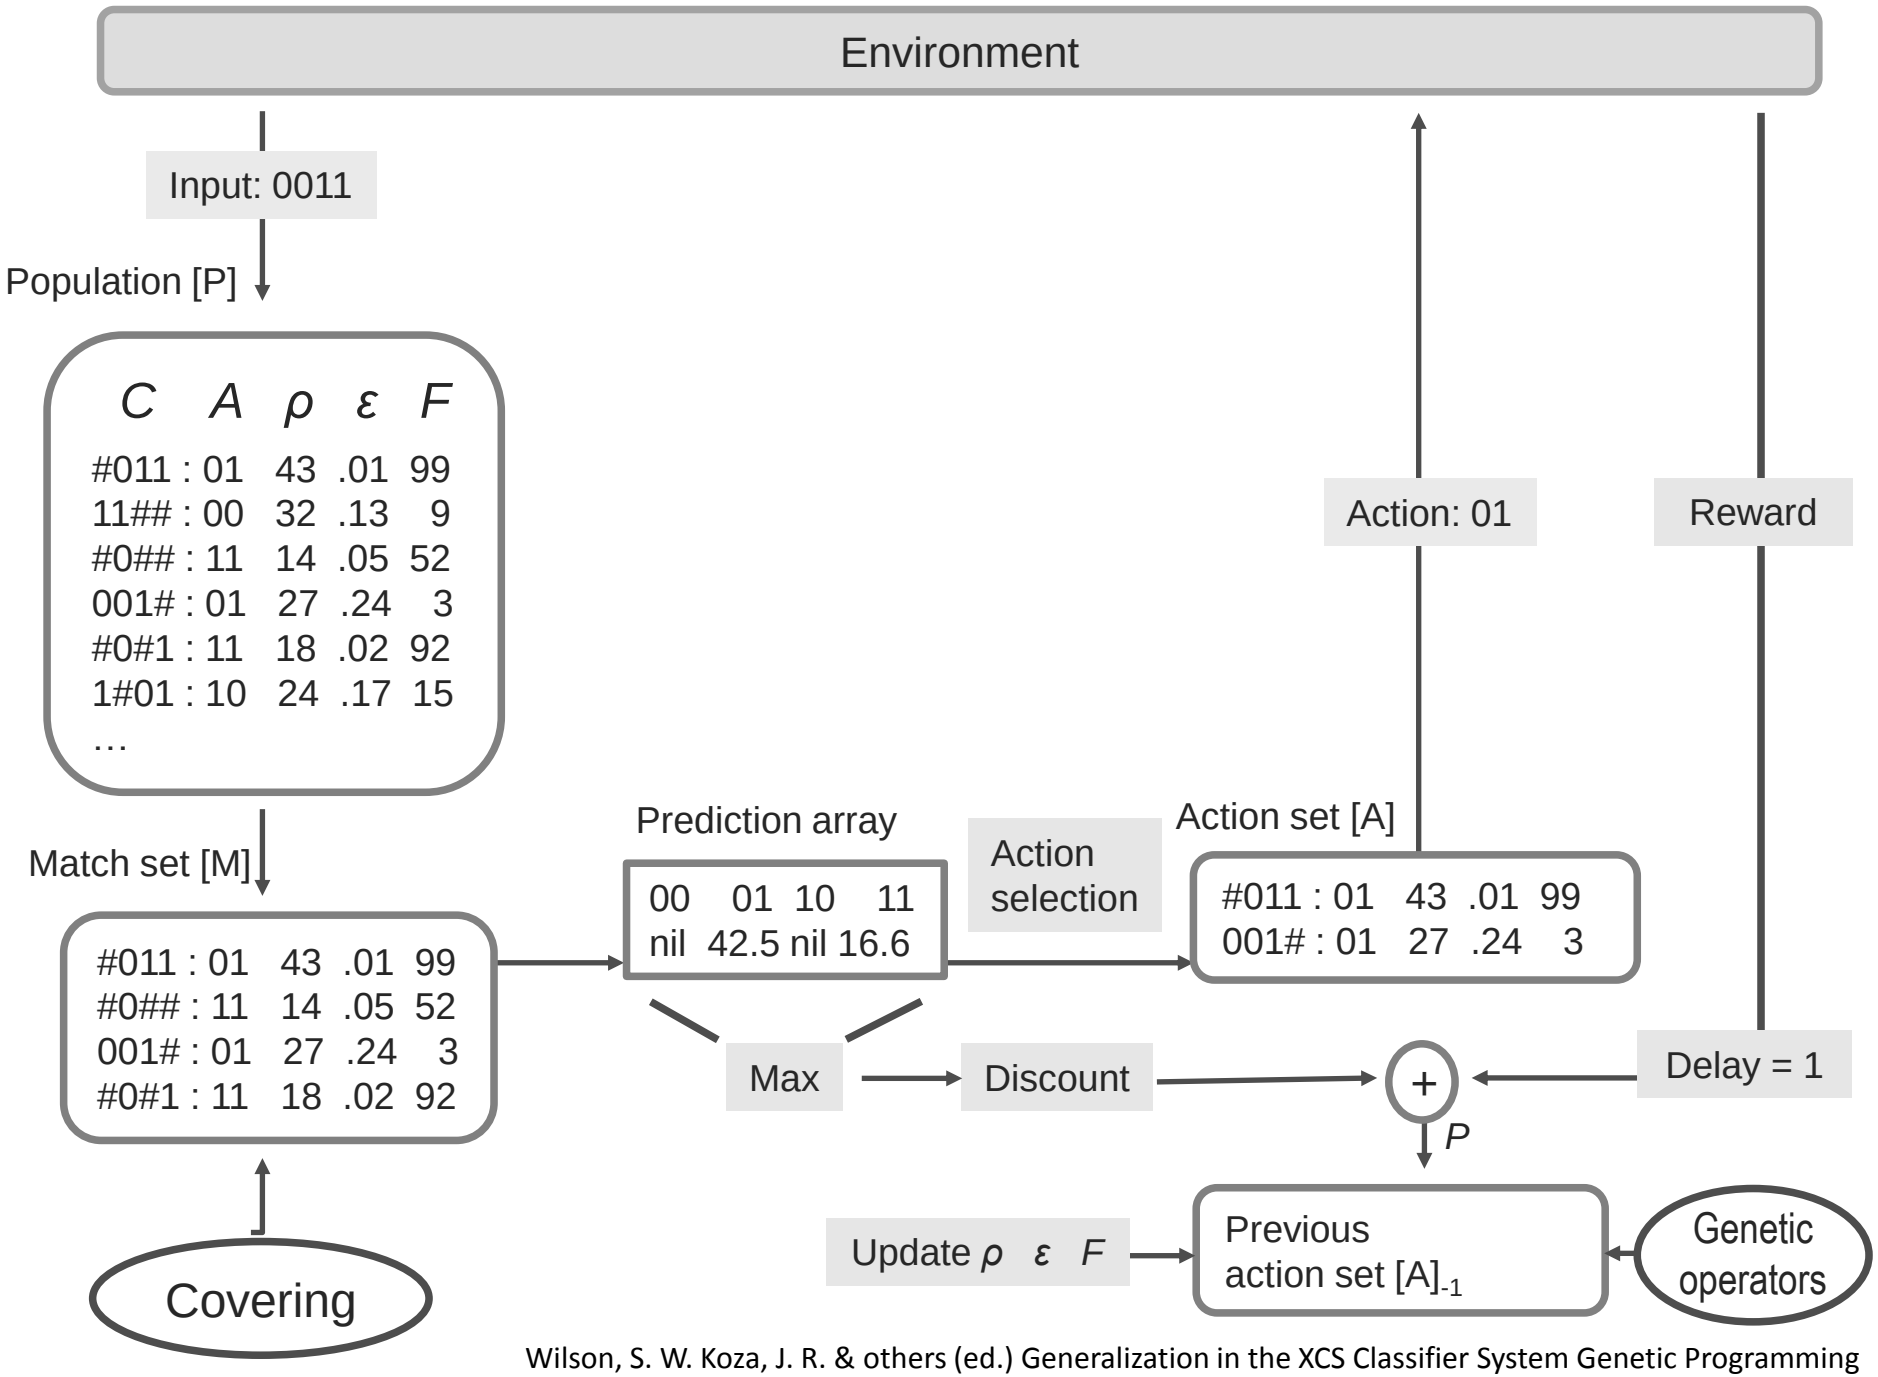

Wilson, S. W. Koza, J. R. & others (ed.) Generalization in the XCS Classifier System Genetic Programming 1998: Proceedings of the Third Annual Conference, 1998, 665-674

Learning Classifier Systems

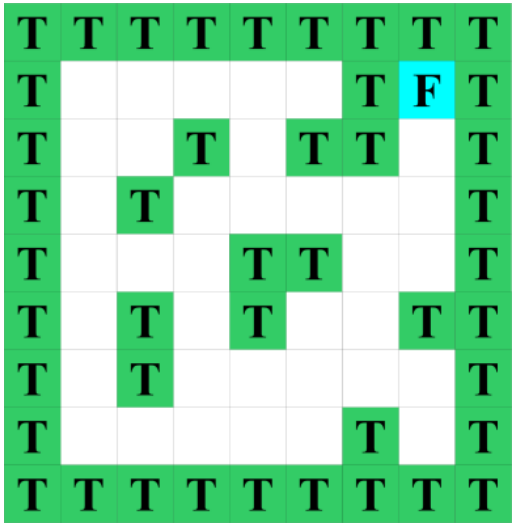

T: Tree
F: Food

- Standard (Multi-Step) Problem:
  - Maze6
- Goal:
  - Find the shortest path to from a random position to food

Learning Classifier Systems

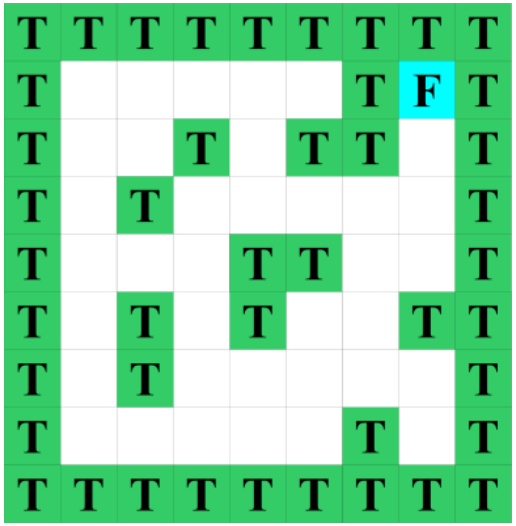

T: Tree
F: Food

- Problem:
  - Limited sensors, no global knowledge
    - Partially observable Markov decision process
- Solution:
  - Iterations, back-propagation of reward

Learning Classifier Systems

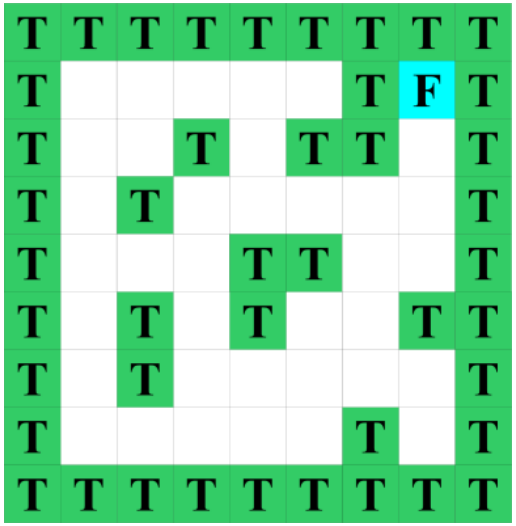

T: Tree
F: Food

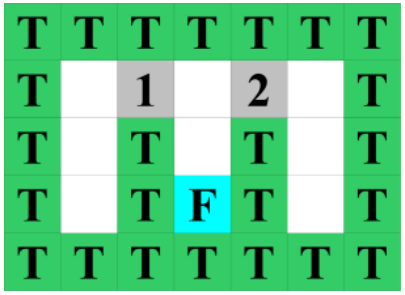

- Problem:
  - Limited sensors, no global knowledge
    - Partially observable Markov decision process
- Solution:
  - Iterations, back-propagation of reward
- Aliasing positions:
  - Handle by using memory

Predator/Prey Scenarios

- Many aliasing positions
- Other agents present
- Dynamic world
  - food and other agents move
- Limited sensors
  - ca. 5 cells range

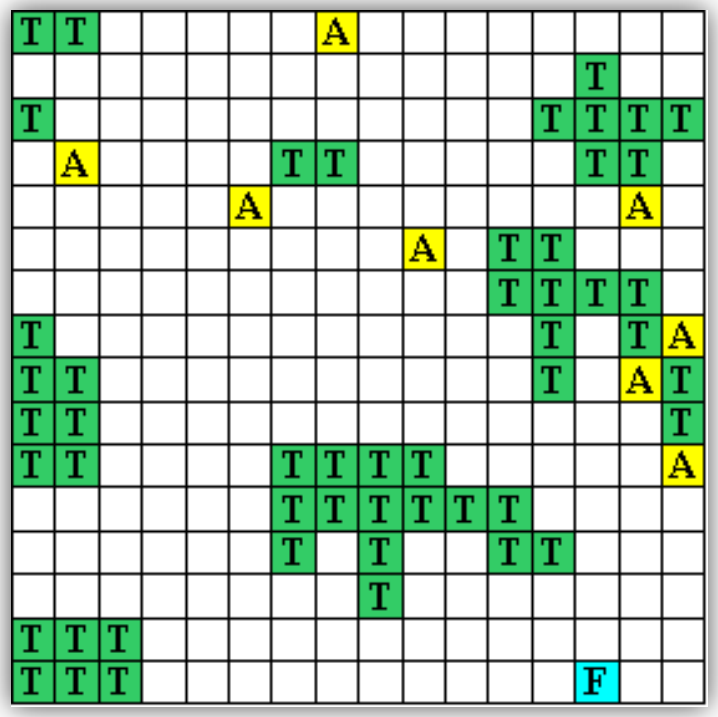

T: Trees
F: Food
A: Agent

Predator/Prey Scenarios

- Terminology:
  - Obstacles, prey, predator
- Goal: Try to stay near the prey
  - Global observation task
  - Runs continuously
  - Maximize average quality

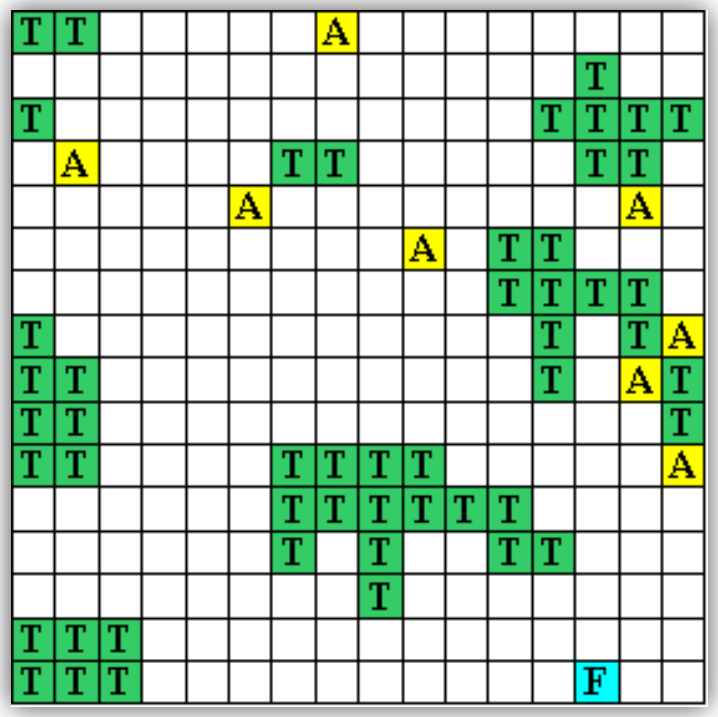

T: Trees/Obstacles
F: Food/Prey
A: Agent/Predator

Classification of Predator/Prey Scenarios

(1) Access to local information only

(2) Open areas with some obstacles

(3) Internal state unknown to others

- No standard MDP
  - Limited sensors (1, 3)
  - Aliasing positions (2)

Classification of Predator/Prey Scenarios

(1) Access to local information only

(2) Open areas with some obstacles

(3) Internal state unknown to others

(4) Dynamic scenario

- No standard MDP
  - Limited sensors (1, 3)
  - Aliasing positions (2)
- No POMDP
  - Non-static scenario (4)

Classification of Predator/Prey Scenarios

(1) Access to local information only

(2) Open areas with some obstacles

(3) Internal state unknown to others

(4) Dynamic scenario

(5) Predators share global observation task

(6) Runs continuously

- No standard MDP
  - Limited sensors (1, 3)
  - Aliasing positions (2)
- No POMDP
  - Non-static scenario (4)
- XCS has to be adapted
  - No "final" reward (5), no iterations (6)

Adaption of the Standard XCS Reward Function

- Standard implementation:
  - Reward:
    - Prey is in a neighboring cell

- Adapted implementation
  - Reward:
    - Prey is in observation range
      - "XCS obs"
    - Prey is in sight range
      - "XCS sight"

Adaption of the Standard XCS Reward Function

- Standard implementation:
  - Reward:
    - Prey is in a neighboring cell
  - Reward action:
    - Assign reward
    - Restart scenario
    - Explore/exploit phase
- Adapted implementation
  - Reward:
    - Prey is in observation range
      - "XCS obs"
    - Prey is in sight range
      - "XCS sight"
  - Reward action:
    - Assign reward
    - Continue scenario
    - Always use exploit phase

Sensors

One sensor array for each direction

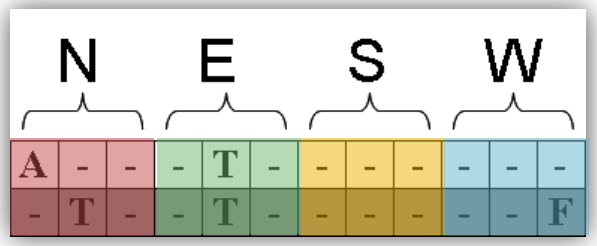

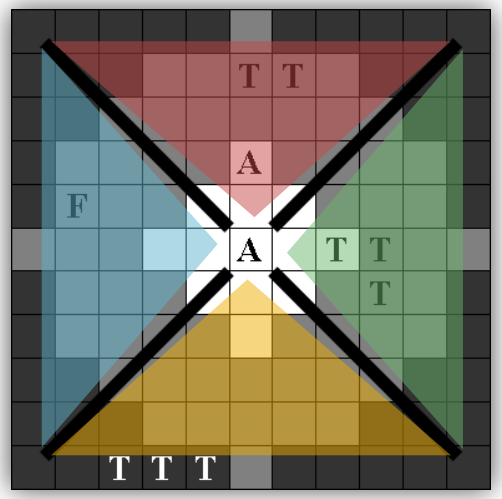

Sensors

One sensor array for each direction

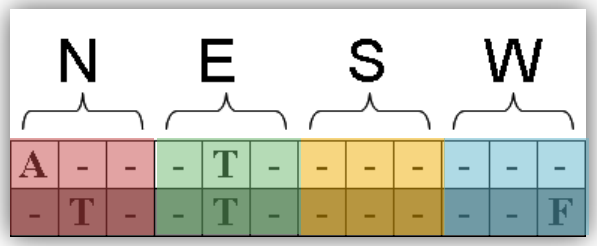

Sensors can sense either far or near (observation range / sight range)

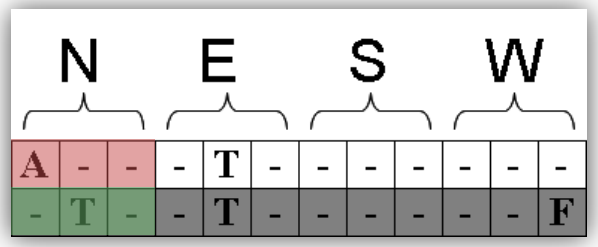

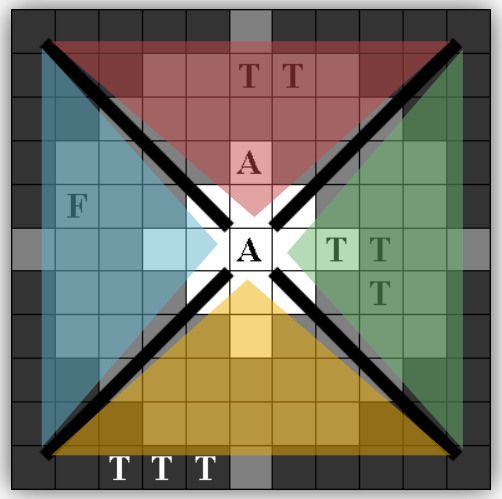

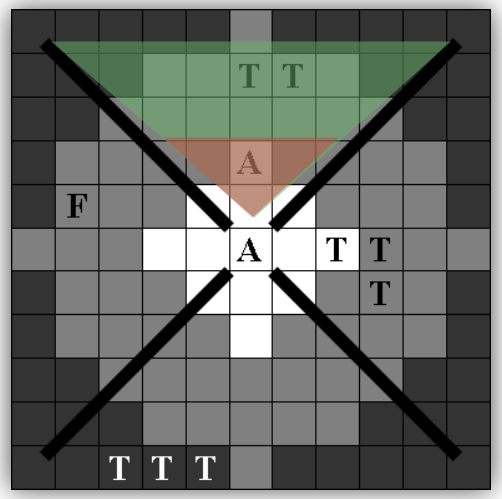

Sensors

One sensor array for each direction

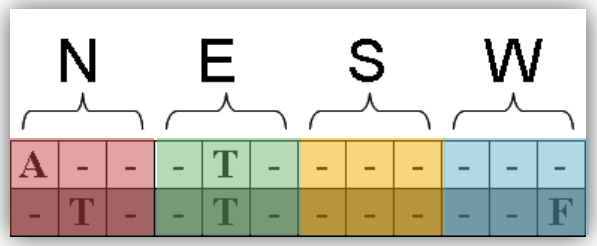

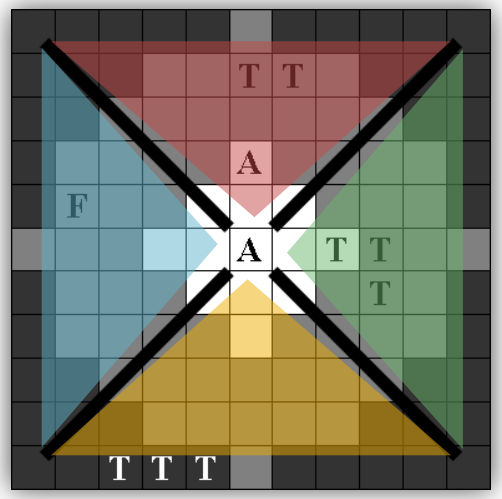

Sensors can sense either far or near (observation range / sight range)

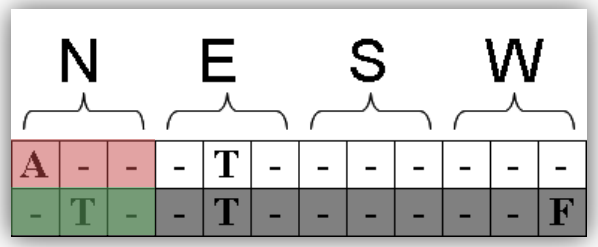

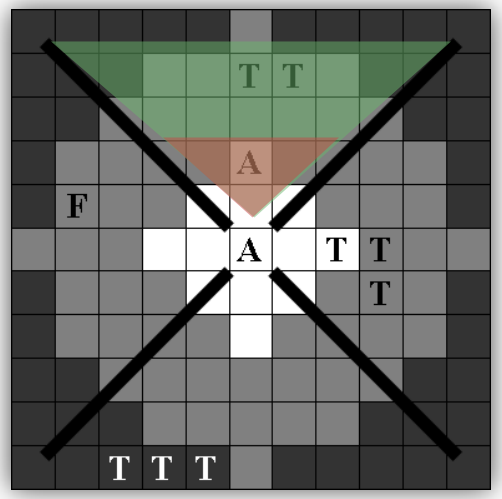

Sensors can distinguish between predators, prey, and obstacles

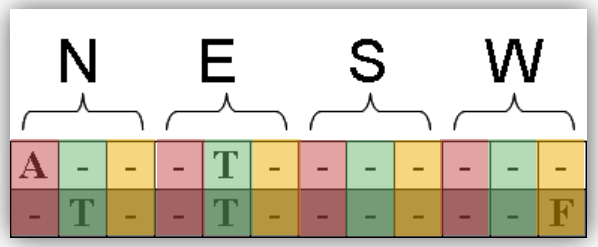

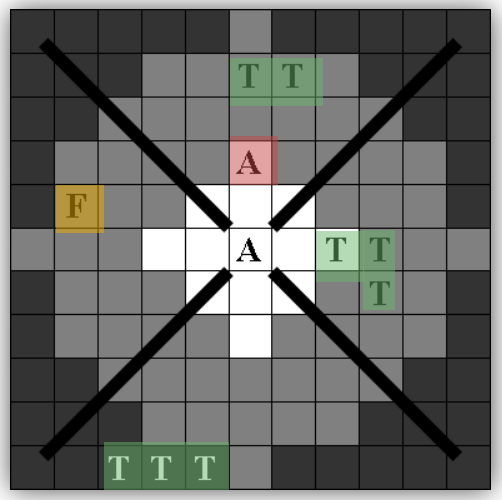

Testing Methodology

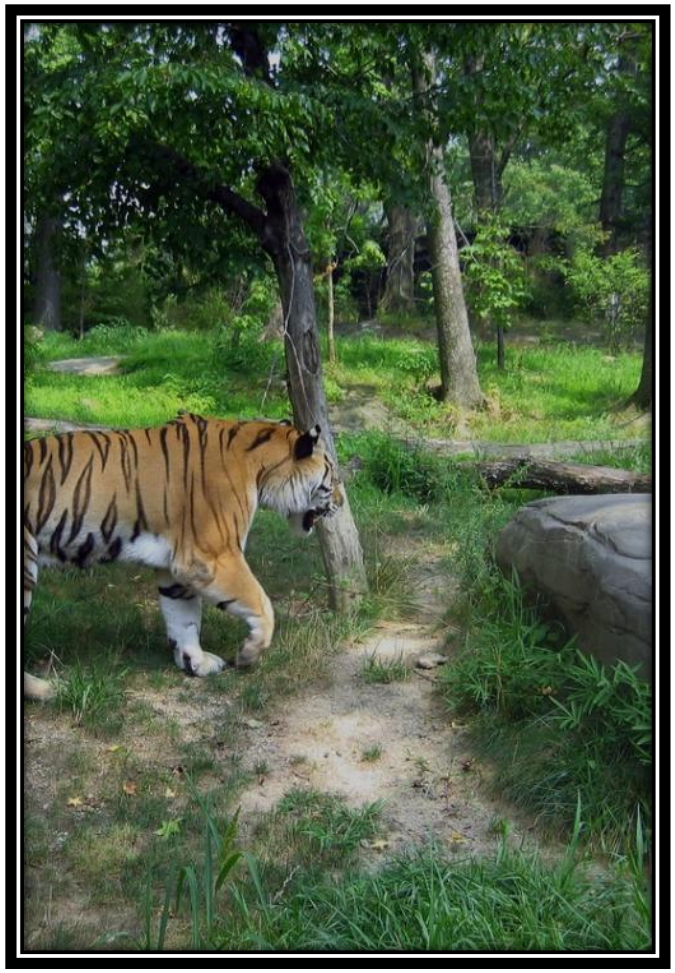

http://www.flickr.com/photos/james_crowley

- Behavior
  - 1 Prey
    - "Obstacle-evading prey"
    - "Predator-evading prey"
    - "Blinded Prey"
    - Speed 2 cells / step
  - 8 Predators
    - Speed 1 cell / step
- Standard XCS parameter settings

Testing Methodology

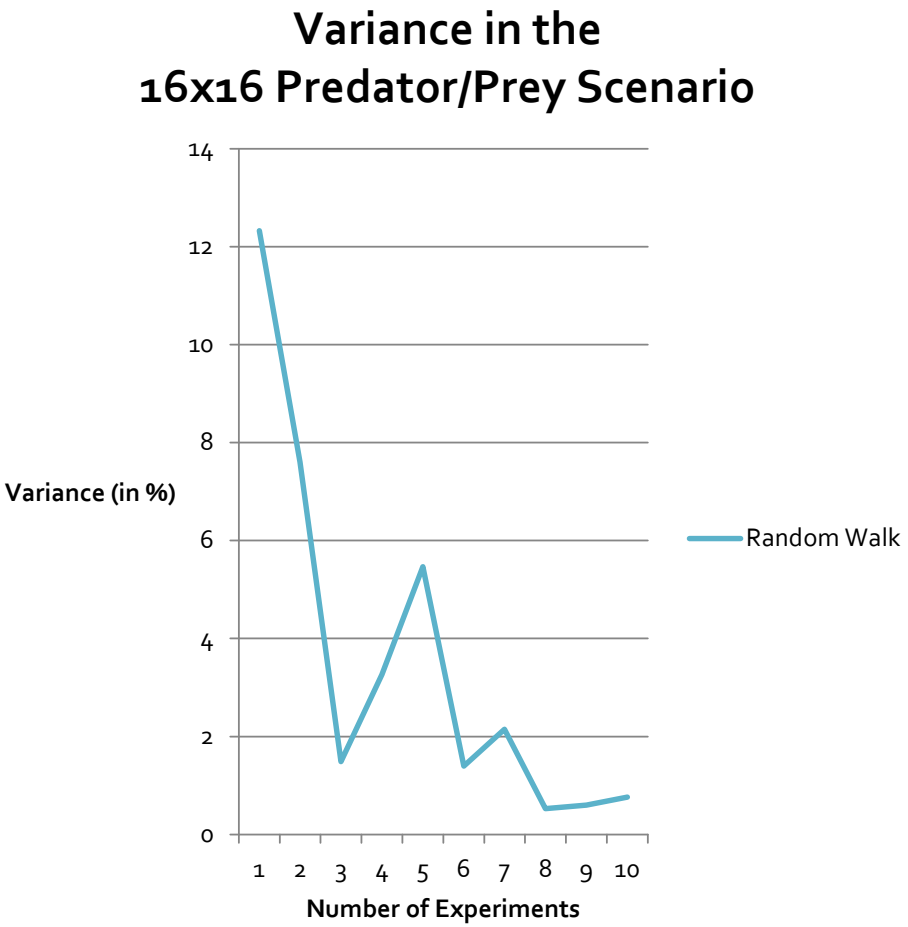

- 2,000,000 steps
- Reset of XCS every 20,000 steps (="experiment")
- Reset of scenario (new random positions) every 2,000 steps

XCS Experimental Results "Pillar Scenario"

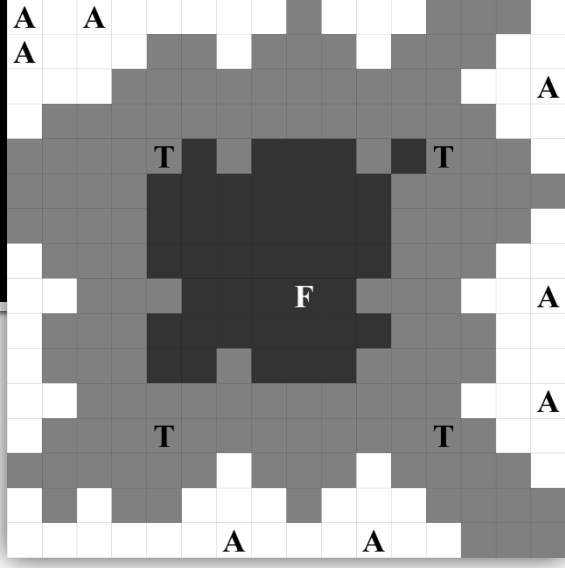

- XCS (obs) shows some learning

Obstacle-evading prey

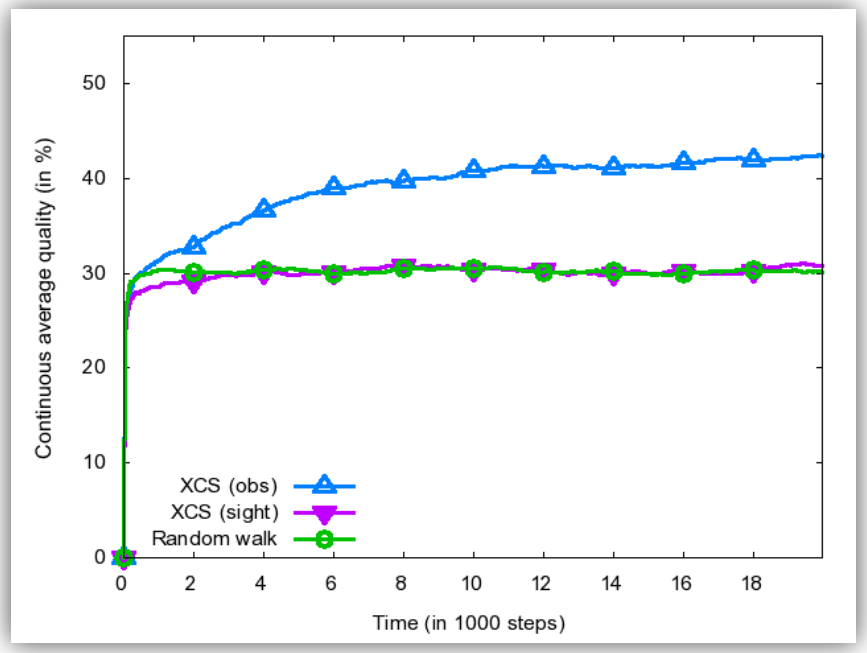

Predator-evading prey

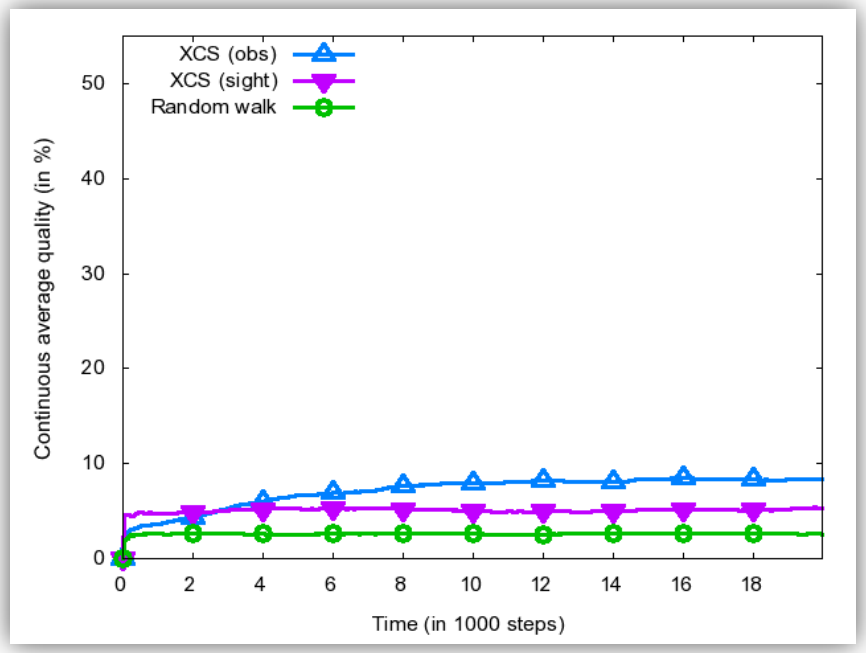

XCS Experimental Results "Random Scenario"

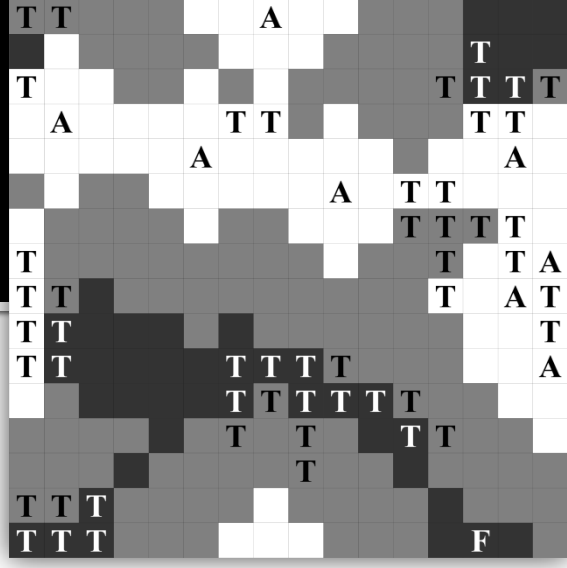

- XCS shows very little learning

Obstacle-evading prey

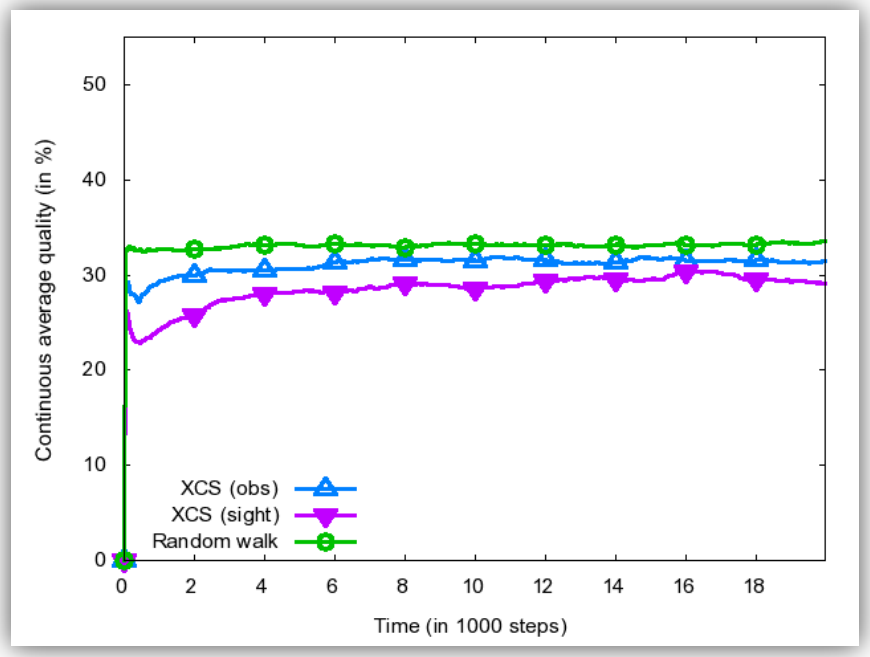

Predator-evading prey

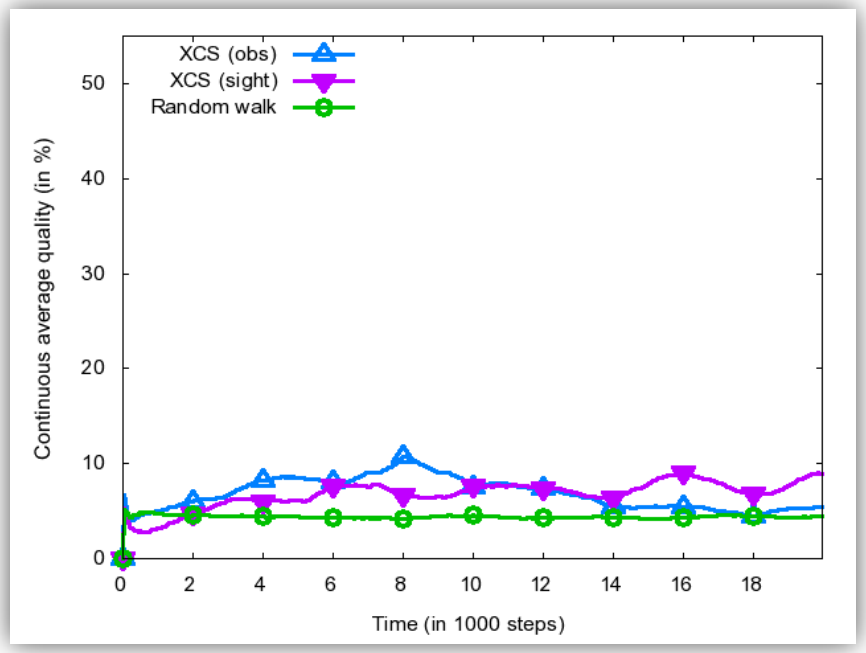

XCS Experimental Results "Difficult Scenario"

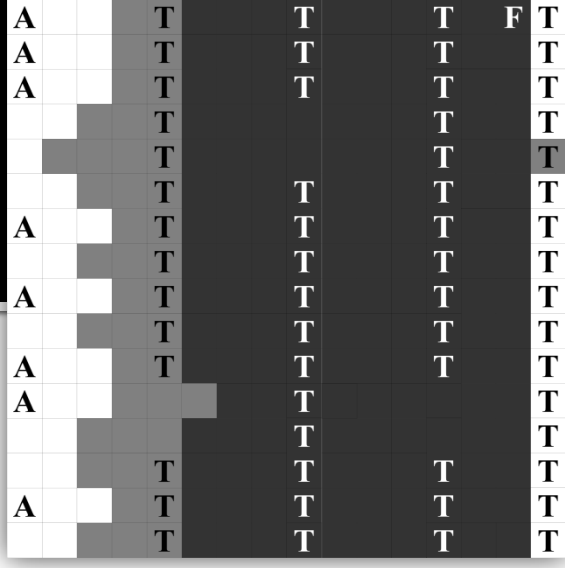

- XCS shows significant learning
  - But also unlearning after 8,000 steps
- "Difficult Scenario" is a maze-like scenario, this result was expected

Blinded prey

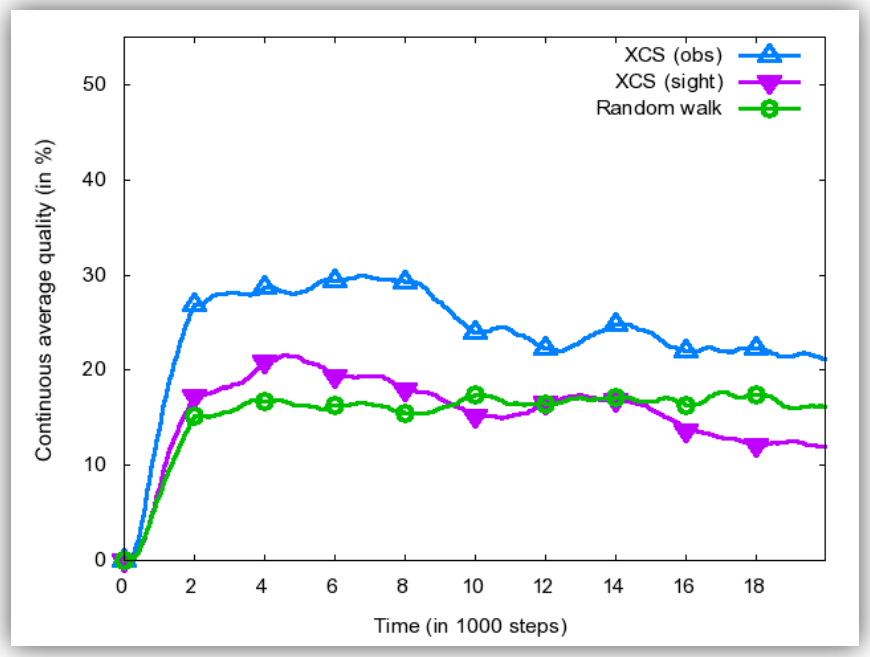

Reward Events "eventXCS"

- Move West

	N	E	S	W
A	1	0	0	0
T	0	1	0	0
F	0	0	0	0
A	0	0	0	0
T	1	1	0	0
F	0	0	0	1

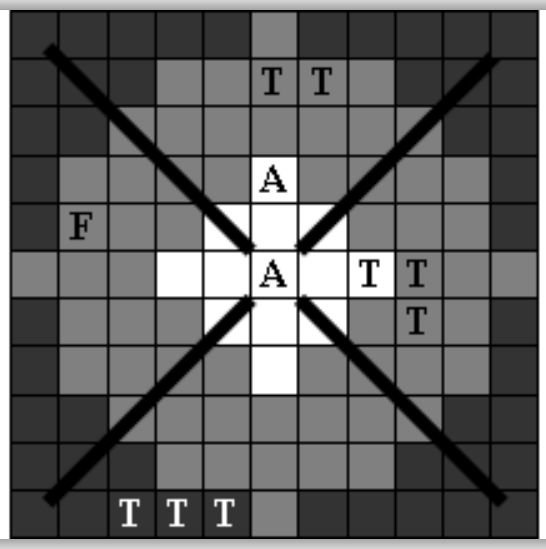

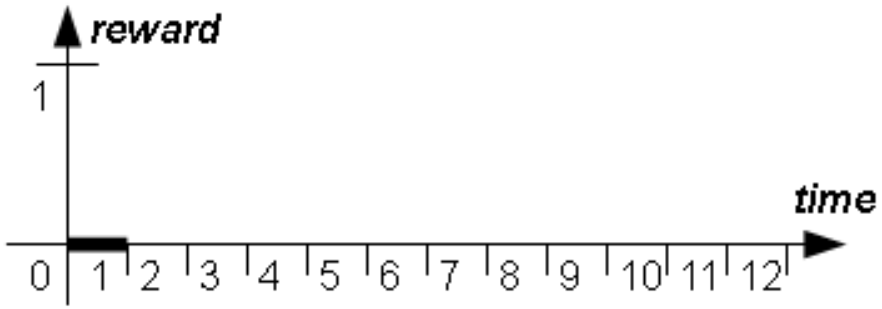

Reward Events "eventXCS"

- Move West

	N	E	S	W
A	0	0	0	0
T	0	0	0	0
F	0	0	0	0
A	1	0	0	0
T	1	1	1	0
F	0	0	0	1

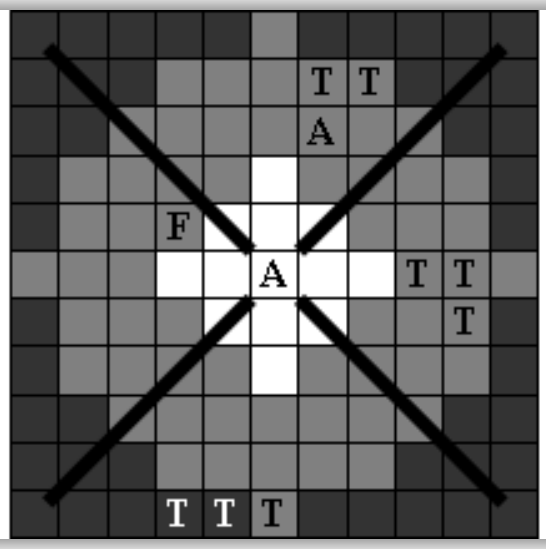

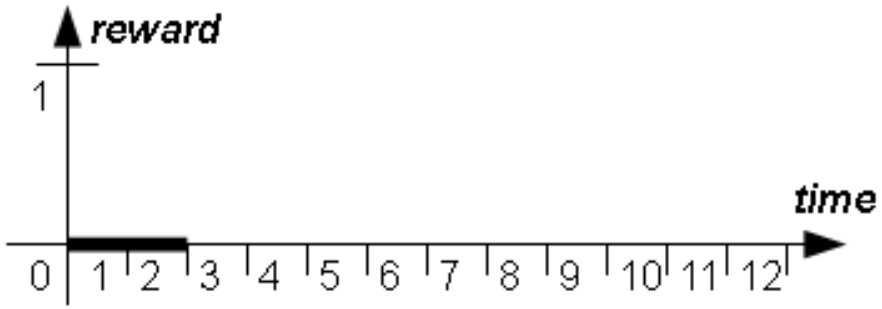

Reward Events "eventXCS"

- Move North

	N	E	S	W
A	0	0	0	0
T	0	0	0	0
F	0	0	0	0
A	1	0	0	0
T	1	1	1	0
F	1	0	0	0

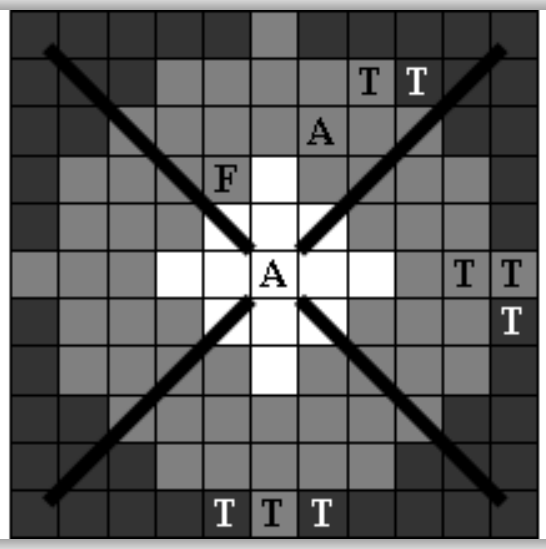

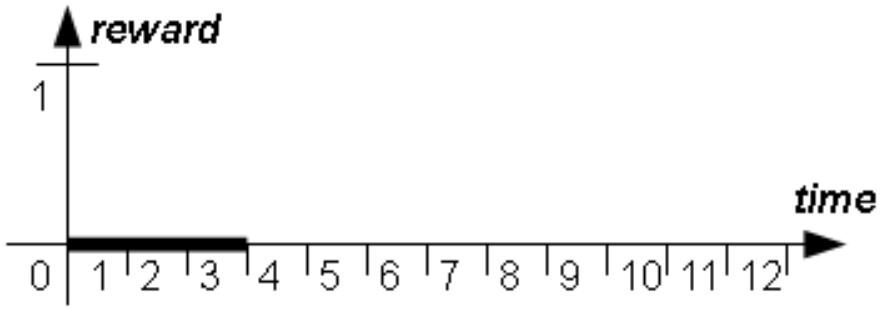

Reward Events "eventXCS"

- Move West

	N	E	S	W
A	1	0	0	0
T	0	0	0	0
F	0	0	0	1
A	0	0	0	0
T	1	1	0	0
F	0	0	0	0

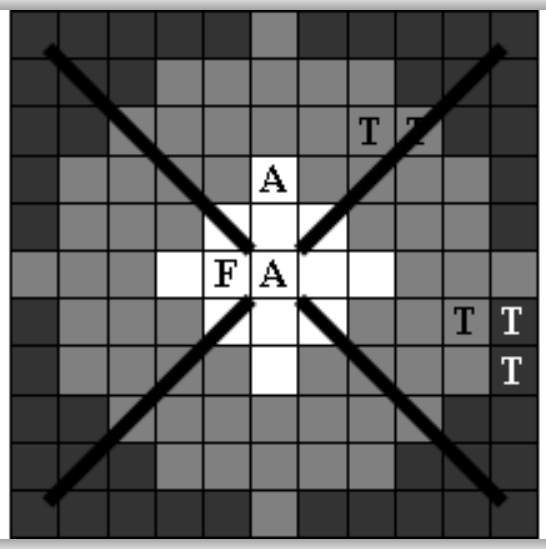

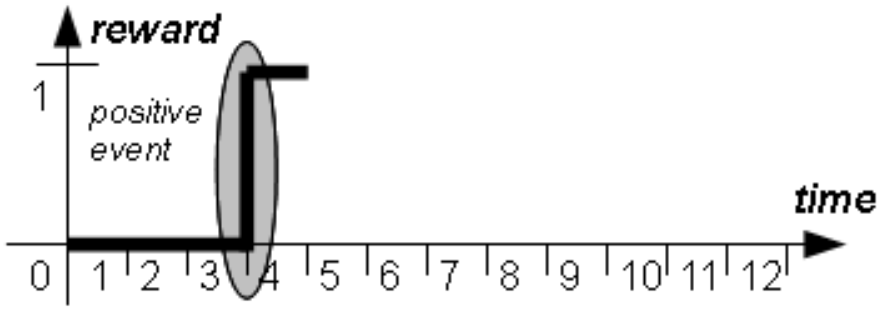

Reward Events "eventXCS"

- Move North

	N	E	S	W
A	0	0	0	0
T	0	0	0	0
F	1	0	0	0
A	1	0	0	0
T	1	0	0	0
F	0	0	0	0

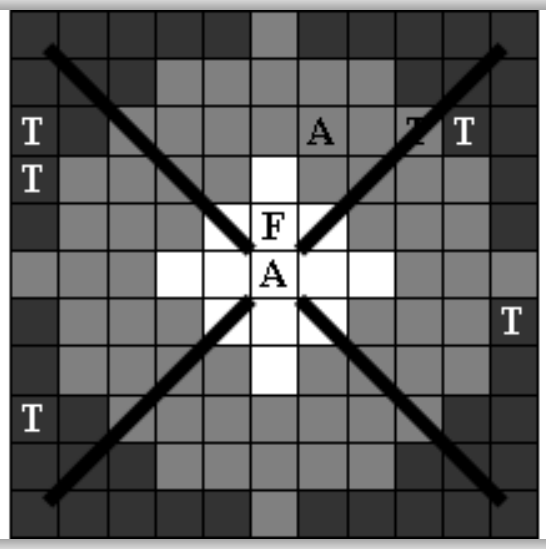

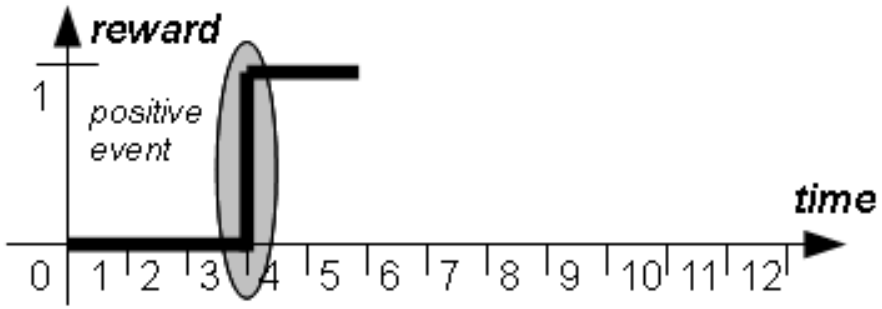

Reward Events "eventXCS"

- Move North

	N	E	S	W
A	0	0	0	0
T	0	0	0	0
F	1	0	0	0
A	1	0	0	0
T	0	1	0	0
F	0	0	0	0

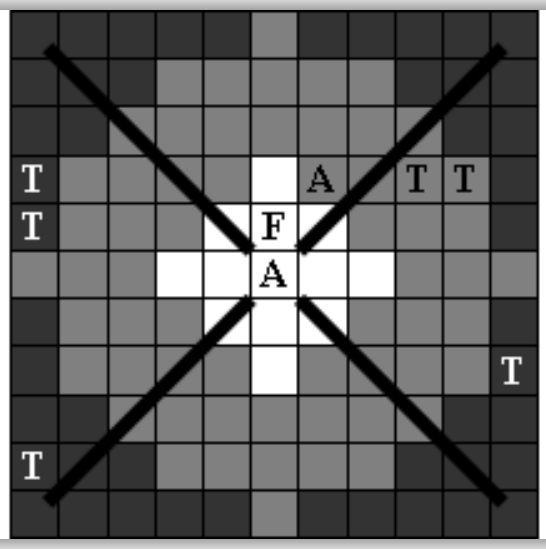

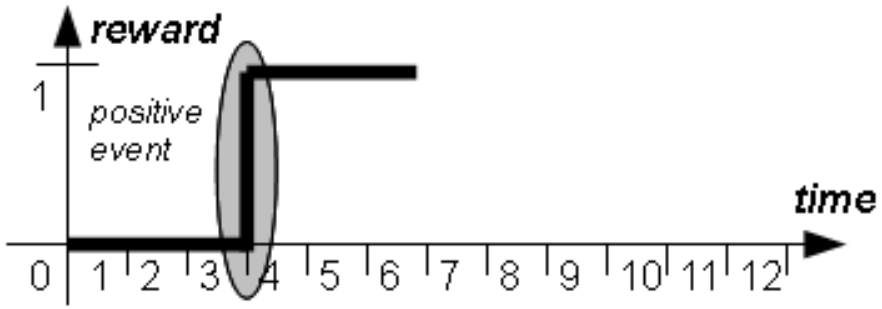

Reward Events "eventXCS"

- Move West

	N	E	S	W
A	0	0	0	0
T	0	0	0	0
F	1	0	0	0
A	1	0	0	0
T	0	1	0	1
F	0	0	0	0

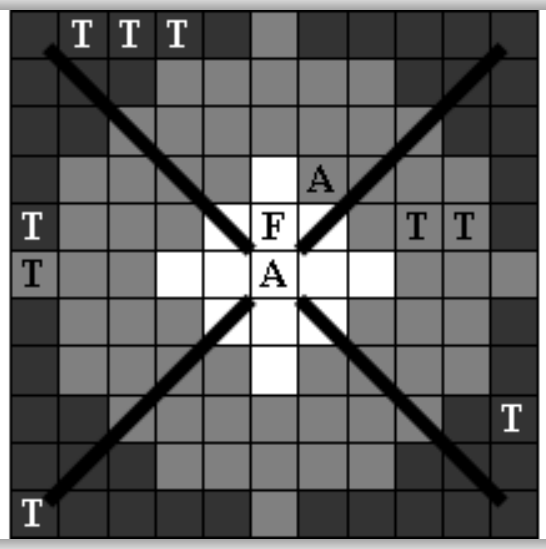

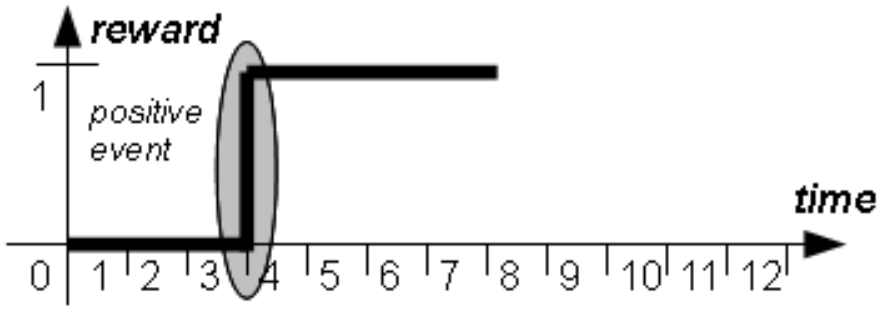

Reward Events "eventXCS"

	N	E	S	W
A	0	0	0	0
T	0	0	0	0
F	0	0	0	0
A	1	0	0	0
T	0	1	0	1
F	0	1	0	0

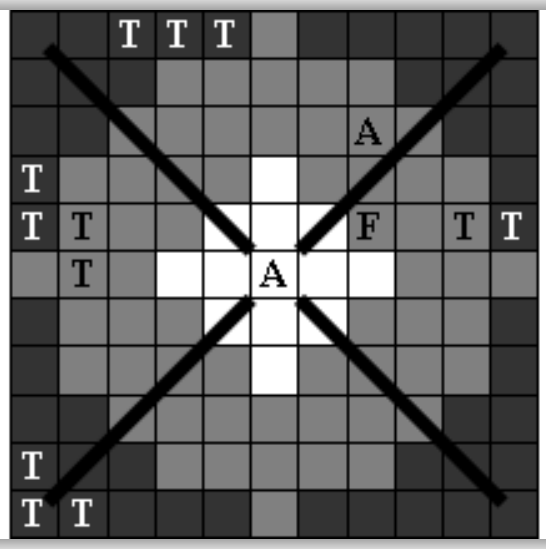

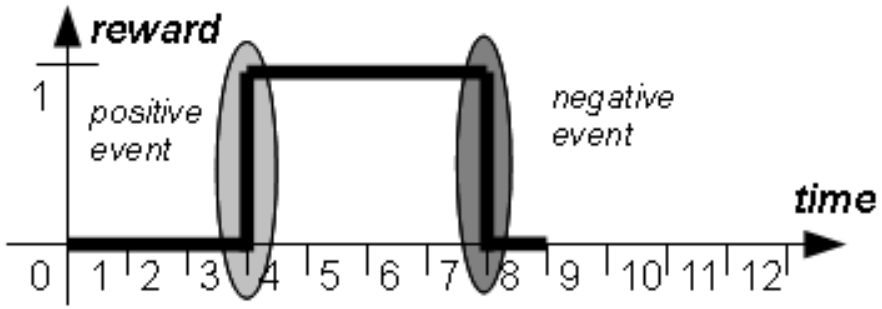

Reward Distribution "eventXCS"

- Analyze succession of positive and negative events

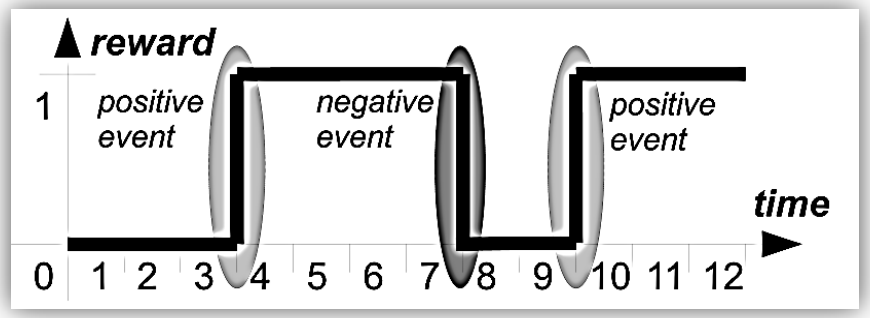

Reward Distribution "eventXCS"

- Analyze succession of positive and negative events
- Distribute the reward as soon as possible (i.e. at each event)
- Idea:
  - Action sets close to an event probably contributed more

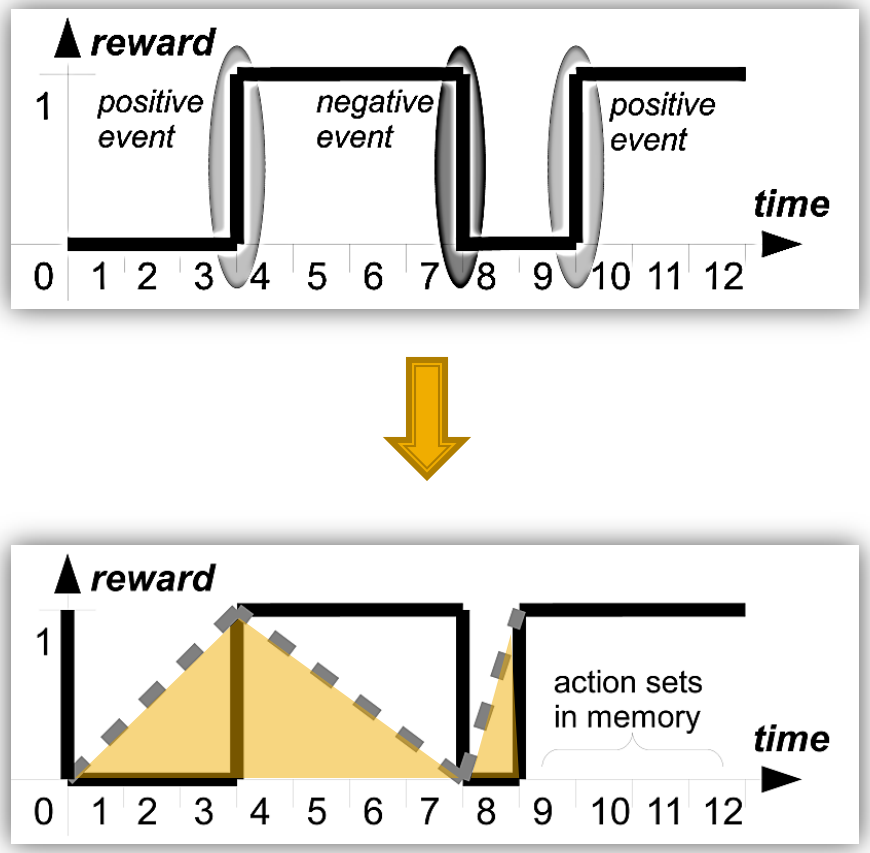

Experimental Results "Pillar Scenario"

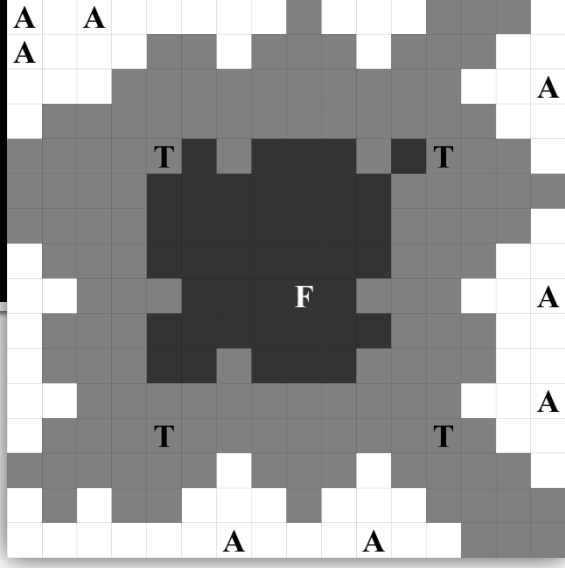

- eventXCS clearly outperforms XCS

Obstacle-evading prey

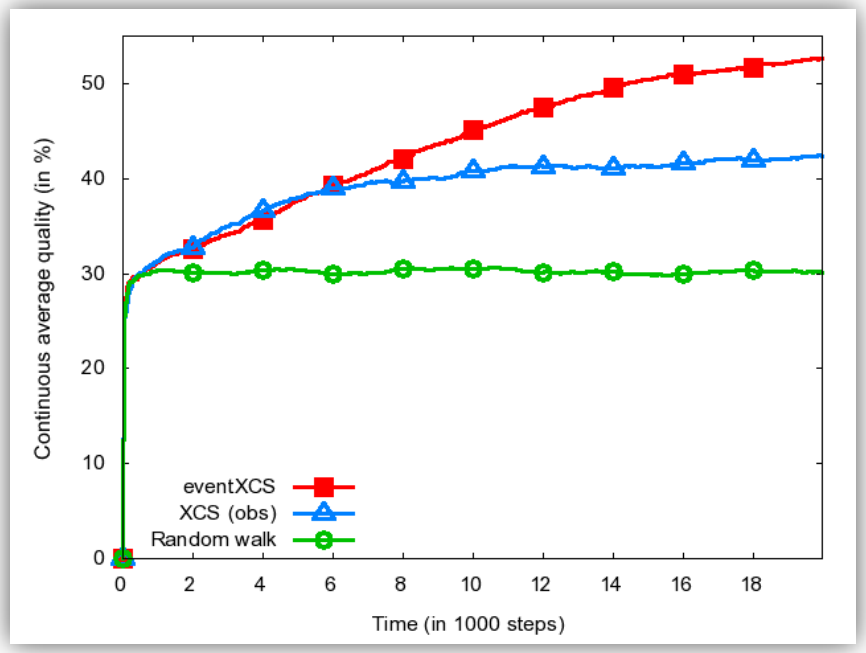

Predator-evading prey

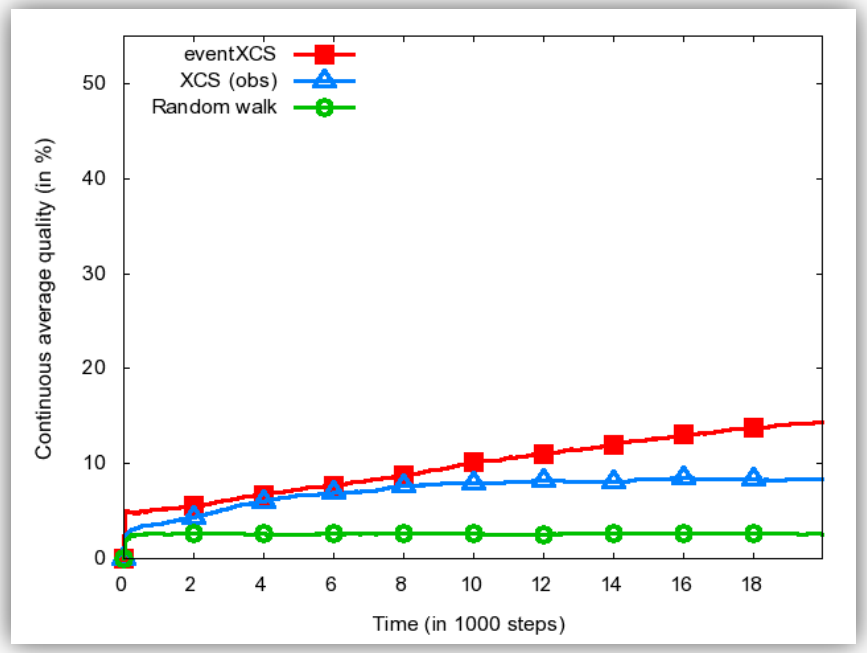

Experimental Results "Random Scenario"

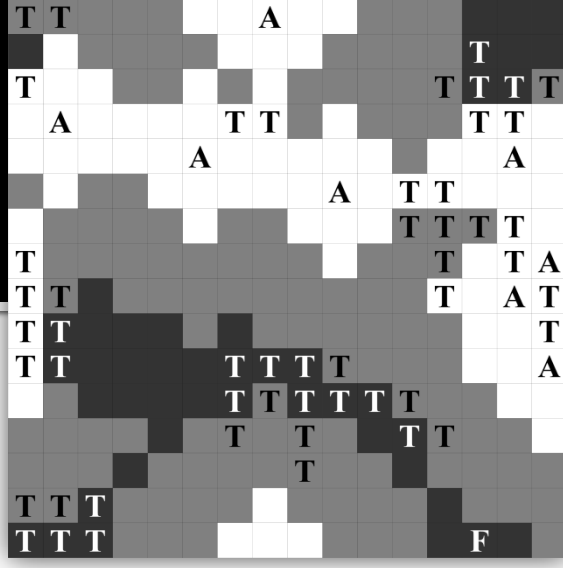

- eventXCS shows slow but steady learning with an obstacle-evading prey

Obstacle-evading prey

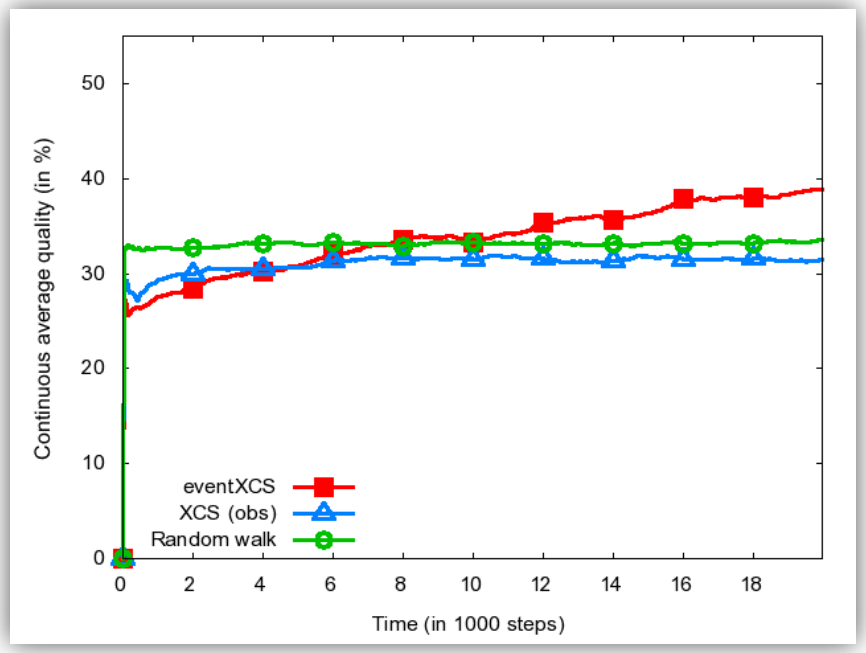

Predator-evading prey

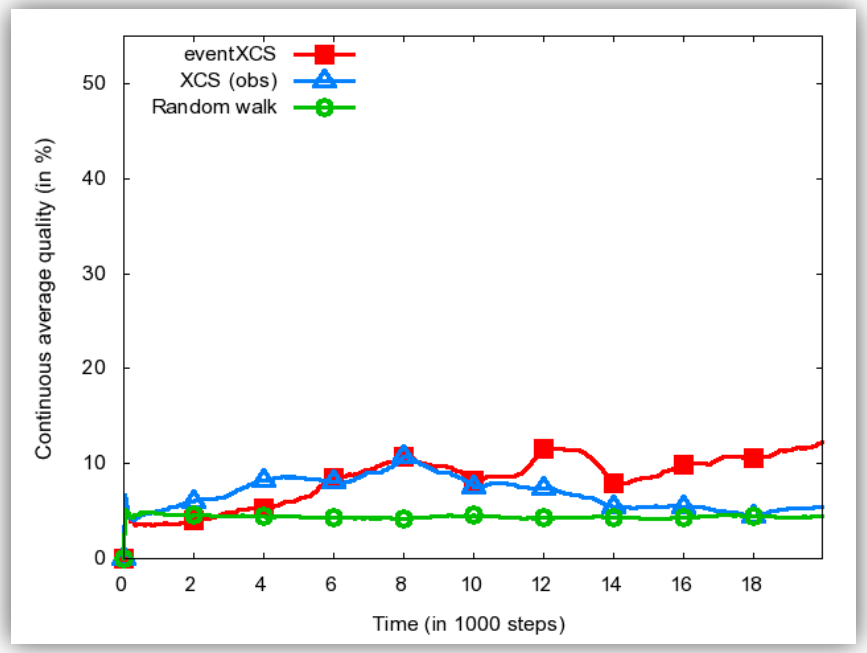

Experimental Results "Difficult Scenario"

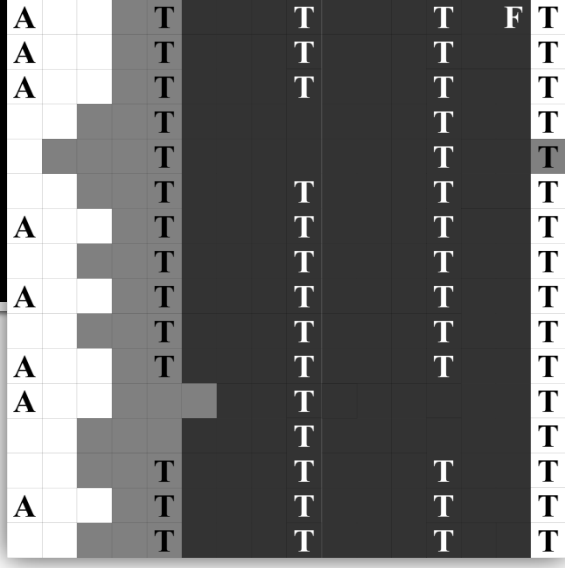

- eventXCS fails in this scenario (not displayed, fitness = ~0)
- Using "tournament selection" shows acceptable results with no sign of unlearning

Blinded prey

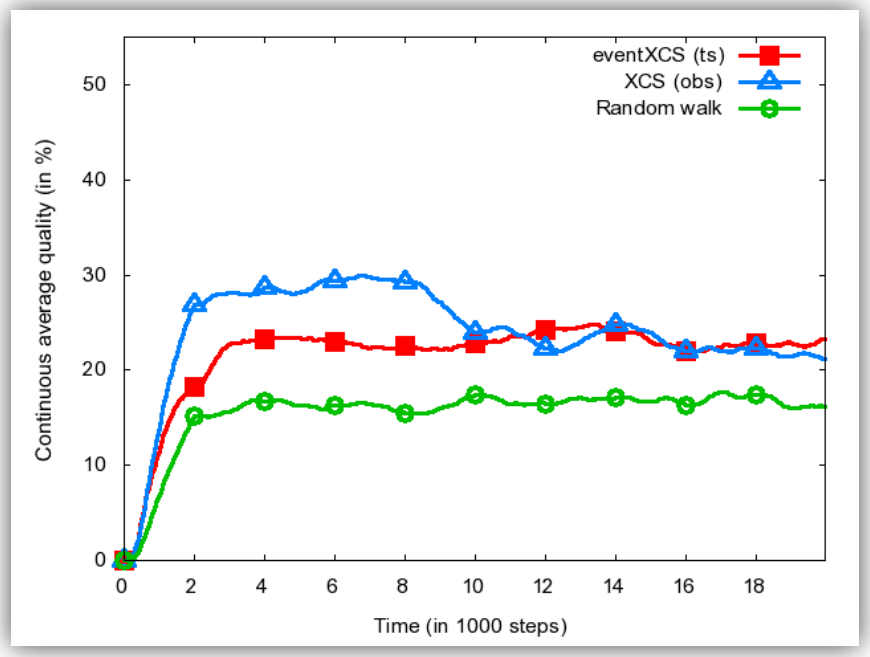

Conclusion

- Predator/Prey scenarios are NOMDP
- XCS can learn (with minimal adaptions) in some P/P scenarios
- Using event handling and reward distribution (eventXCS) much better learning can be observed
- But: Might need some improvement in difficult scenarios

Backup slides

Neutral Events

- Neutral Event
  - No positive or negative event for a number of steps
  - Half of the action sets is discarded and receives reward
  - Idea:
    - Good actions are rewarded earlier
    - Preventing of dead ends

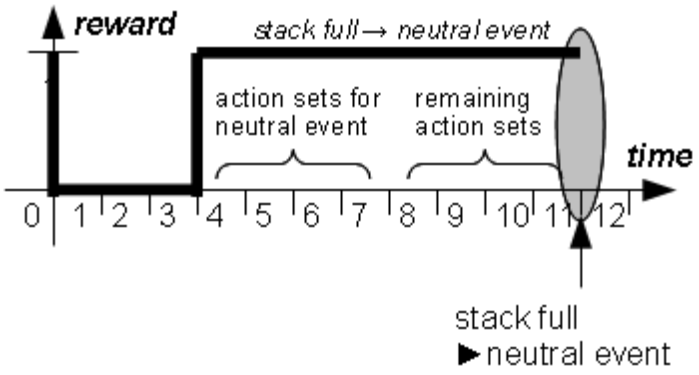

Neutral Events

- Neutral Event
  - No positive or negative event for a number of steps
  - Half of the action sets is discarded and receives reward
  - Idea:
    - Good actions are rewarded earlier
    - Preventing of dead ends
  - Problem:
    - Error possibility high if directly followed by an event.

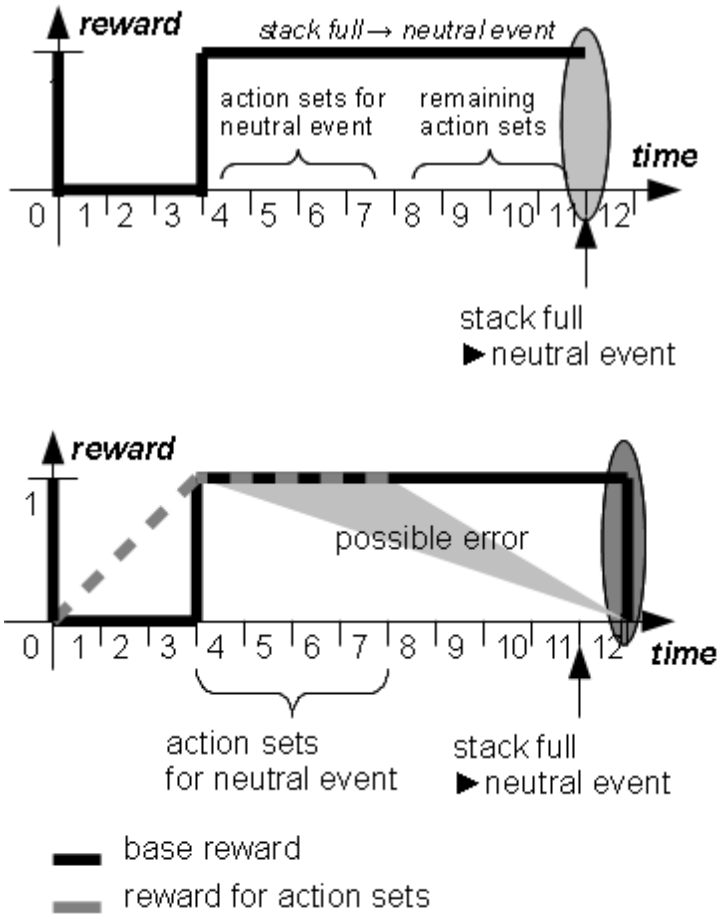

Neutral Events

- Tests have shown that a stack size of 8 is generally good for all three scenarios

"Pillar Scenario"

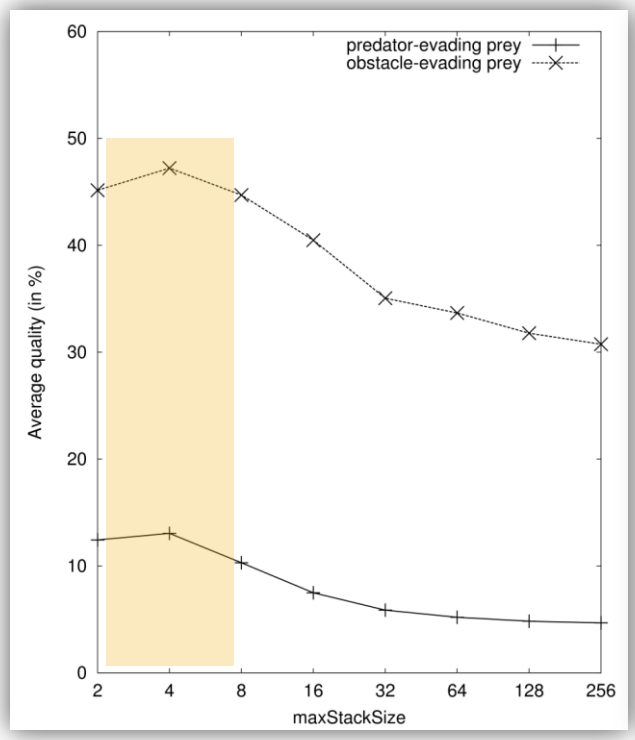

"Random Scenario"

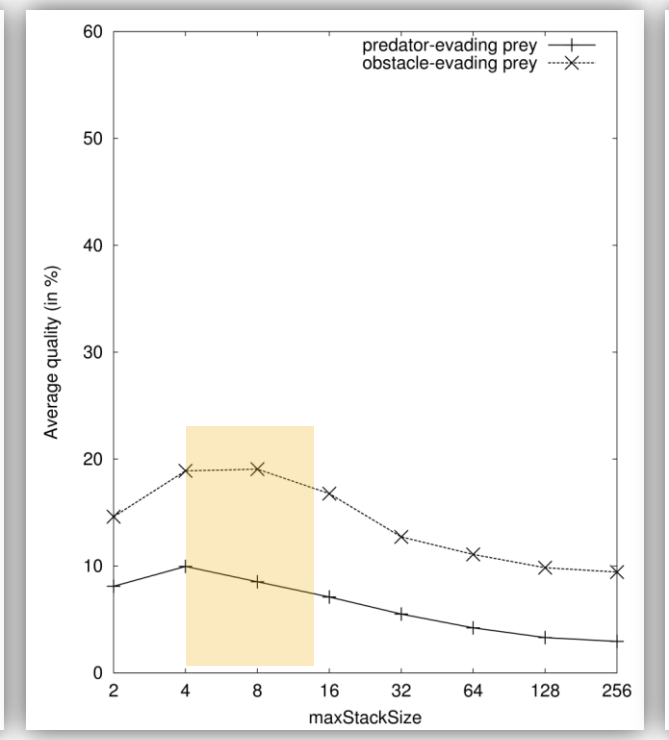

"Difficult Scenario"

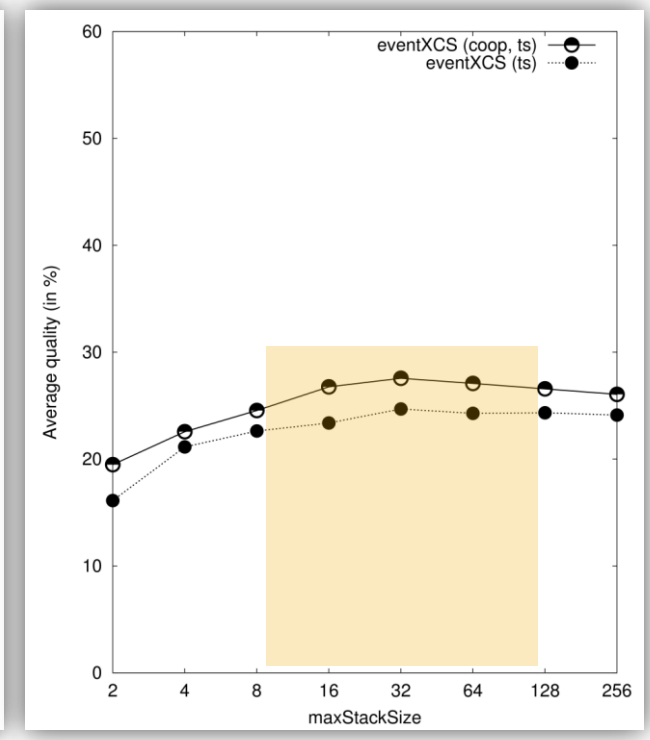

Learning Rate β

- Pillar Scenario
  - Obstacle-evading prey
  - Low learning rate (0.05) good, eventXCS very stable

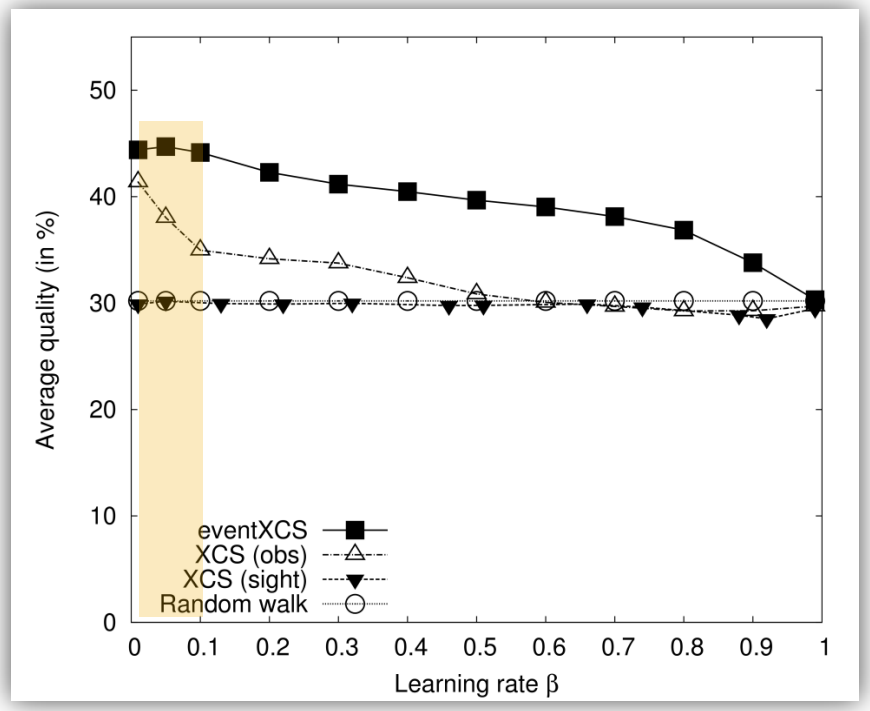

Learning Rate β

Pillar Scenario
Predator-evading prey

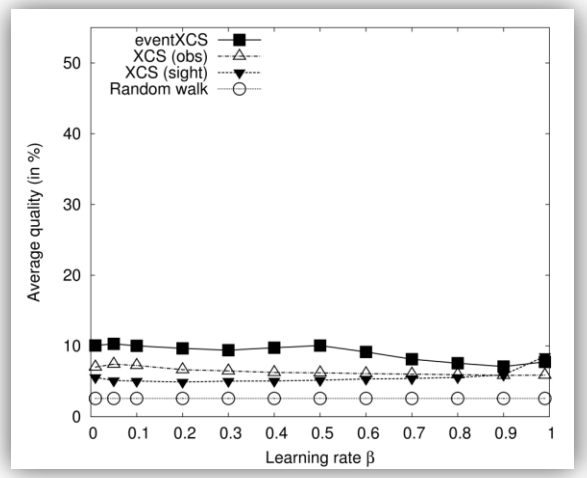

Random Scenario,
Obstacle evading prey

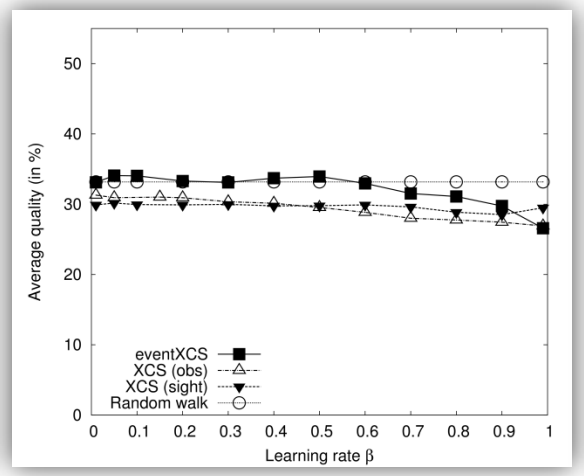

Random Scenario,
Predator evading

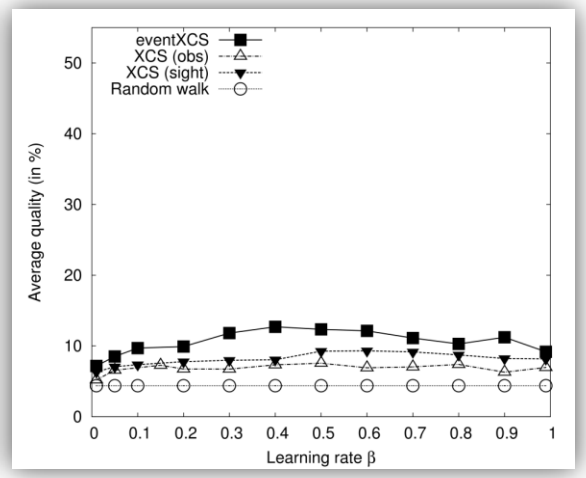

Learning Rate β

- Difficult Scenario
  - Blind prey
  - High learning rates show an advantage because of long distance to the prey